

Foundational Principles of Binary Mathematics: A Machine Epistemology

Abstract: This chapter introduces *Binary Mathematics* not merely as a collection of structures, but as a philosophical and computational framework grounded in **Machine Epistemology**. We articulate its core epistemological and ontological principles, positing that all verifiable mathematical knowledge arises from structures ultimately representable and executable via finite operations on a fundamental binary substrate. We argue that mathematics, viewed through this lens, constitutes a hierarchy of *Domain-Specific Languages (DSLs)* facilitating structured symbolic compression over binary distinctions, thereby bridging abstract reasoning with machine-resolvable verification. The focus here is on the **why** and **what** of this foundation, deferring the detailed computational **how** to subsequent chapters. All concepts of truth, proof, number, set, and function are re-examined in terms of executable binary state transformations, laying a rigorous groundwork for an **executable, verifiable mathematics** aligned with the principles of Machine Epistemology.

1.1 Philosophical Positioning: Machine Epistemology as Fundamental Stance

We begin by establishing **Machine Epistemology** as our philosophical baseline. This perspective prioritizes the conditions under which mathematical knowledge can be represented, verified, and utilized by information processing systems (whether biological or artificial), over traditional debates about the existence of mathematical objects in a Platonic realm or their conception in human intuition. In contrast to viewing mathematics as a purely human linguistic construct or a product of intuition, Machine Epistemology asks: *Under what conditions can knowledge be rendered executable and verifiable by a machine?* The focus shifts from subjective conviction to objective computability and verifiability. In other words, what matters is not whether an argument **convinces** a person, but whether a proposed structure or transformation **computes correctly and verifiably** on a machine.

Operational Duality (Discovery vs. Construction): A key tenet of this stance is an *operational duality* between what is *discovered* and what is *constructed* in mathematics. On the one hand, we **discover** the fundamental logical necessity for distinguishable states in any information-processing system. This necessity gives rise to the most elementary ontological primitive: the binary distinction, 0 vs. 1. In the physical world, as John Wheeler famously put it, “*every it — every particle, every field of force... derives its function, its meaning, its very existence entirely... from bits*”, i.e. binary yes-or-no indications. The binary digit (bit) is the minimal unit of information required to distinguish two equally likely possibilities ¹. This insight, dating back to the foundations of information theory and early computing, is backed by both theory and engineering: Shannon showed that Boolean algebra (true/false, or 0/1) is sufficient to represent any logic circuit or arithmetic operation, and the Church-Turing thesis asserts that any effective computation can be performed with a simple binary-based machine model. In sum, the **binary distinction** ($0 \neq 1$) is *discovered* as an unavoidable foundation for representation and computation – it is the *bedrock of information* and the physical basis of memory in computing systems.

On the other hand, upon this necessary substrate, all higher-level mathematical systems are **constructed**. Humans (or AI agents) invent formalisms – e.g. arithmetic, algebra, set theory, geometry – as structured symbolic frameworks. We term each such framework a *Domain-Specific Language (DSL)* in the broad sense. Each DSL has its own symbols, syntax, and inference rules, actively designed to model some aspect of reality or to solve certain classes of problems. This view aligns with the notion of mathematics as a human-created toolkit (as in formalism), but crucially, it **anchors** every DSL in the binary substrate. There is no contradiction between discovery and invention here: the binary substrate is discovered, while the rich variety of mathematical theories is invented and built atop it (much like software is written to run on binary hardware). The *epistemic consequence* of this duality is that mathematical knowledge is characterized not by platonic truth or communal consensus, but by the ability to **construct, manipulate, and verify** symbolic representations within rule-bound systems that are ultimately traceable to binary operations. In this light, a statement's *truth* is not an ethereal correspondence to an ideal world; rather, truth means that the statement can be represented in a given DSL and that an **execution** of that representation (on appropriate inputs or premises) yields a verifiable result consistent with the statement.

Mathematics as Executable Representation: We thus redefine mathematics as the discipline of creating *precise, executable models of structural relationships*. Mathematical objects like numbers, sets, functions, or geometries are not primarily mystical abstractions or intuitions; they are components of *constructed representational systems* (our DSLs) with explicitly defined roles and relations. Their meaning comes from how they behave within those systems and, ultimately, how they reduce to operations on the binary substrate. This is a shift from traditional views. In classical philosophy of mathematics, one often asks: *Are mathematical entities discovered (Platonism) or invented (Formalism/Intuitionism)?* Here, we answer: the **capacity for mathematics** is discovered (since it requires an information substrate), but the **content of mathematics** is invented (as formal systems built atop that substrate). Importantly, even invented content must **prove itself** by being implementable and verifiable on the substrate. Mathematics, in our view, becomes a *machine-readable, machine-executable knowledge system*. It bridges discernible structural patterns in the world and their verifiable encoding as binary information processes.

Machine Epistemology vs. Traditional Proof: From this machine-centered standpoint, we fundamentally alter the concept of what it means to *know* something in mathematics. Traditionally, mathematical knowledge has been communicated and validated through proofs written in natural or formalized language, intended to persuade human intuition and reasoning. A proof in the traditional sense is a sequence of logical arguments, each justified by axioms or inference rules, convincing a (human) reader that a theorem is true. However, as the QED project noted, fully formal proofs in practice end up looking much like programs – a series of instructions that a computer (or a human following mechanical rules) could check ². Machine Epistemology embraces this reality: we treat proofs *as programs* and mathematical propositions *as executables*. This is reminiscent of the Curry–Howard correspondence in logic and computer science, which establishes a deep analogy between proofs and programs (and between logical propositions and types in a programming language) ³. Indeed, if a proof is a kind of program, then *checking* a proof's correctness is akin to running a program to see if it produces the expected outcome – a fundamentally mechanical process.

In our framework, then, **truth** is defined operationally. A mathematical statement encoded in a particular DSL is considered *true (within that DSL)* if and only if its corresponding executable representation can terminate in a state that satisfies the statement's conditions according to the rules of that DSL. In other words, truth is *that which a correctly constructed computation can demonstrate*. This orientation shifts the locus of certainty from human intuition to machine verification. A complex theorem, no matter how

unintuitive, is accepted as true if we have a verifiable *computation* (proof script, program, or derivation) that produces it from accepted premises. This is aligned with practices in formal verification and theorem proving: for example, the Feit-Thompson theorem (a deep result in group theory) was proven with a six-year effort in the Coq proof assistant, yielding a fully machine-checkable proof ⁴. The computer's confirmation in such cases supplants the need for human consensus; if the formal proof checks out, the result is considered established. We extend this principle universally: all mathematical truths should, in principle, be checkable by executing an appropriate binary program. By tying truth to **successful computation**, we also inject a notion of thermodynamic or information-theoretic efficiency – the execution should lead to a *stable, low-entropy state* that embodies the truth claimed (a notion that will gain significance when we discuss optimization and learning in later chapters).

Finally, **verification** in Machine Epistemology is multi-faceted and machine-centric. It consists of (a) executing the representation (does it run to completion without error?), (b) rule-checking each step against the DSL's formal semantics (does each transformation conform to allowed rules?), and (c) optionally linking or auditing the process with external means like cryptographic hashes or cross-verification with independent systems. This approach has parallels to **Proof-Carrying Code (PCC)** in computer security, where code comes accompanied by a machine-checkable proof that the code respects certain safety properties. In PCC, the onus is on the code producer to provide a valid proof, and the code consumer (verifier) simply checks it quickly. Similarly, in our mathematical context, a proof's author provides a fully detailed executable proof artifact, and any verifier (human or machine) just re-runs and checks it. The result is that mathematical knowledge becomes **self-certifying**, much like PCC ensures programs are *intrinsically* safe by virtue of their attached proofs. There is no need to trust the reputation of a mathematician or the opinion of experts; the proof can be independently validated by anyone with the appropriate verifier (e.g. a proof assistant or our proposed Binary Decomposition Interface execution environment).

In summary, Machine Epistemology reframes mathematics from a system of convincing arguments to a system of verifiable processes. It asserts: (1) The only *primitive* we truly require is the binary distinction, discovered as the basis of information. (2) All else in mathematics is a *constructed language* (DSL) that must prove its worth via executable instantiation. (3) Mathematical truth is an emergent property of successful computations, not of metaphysical alignment with an ideal world. (4) Proofs are not rhetorical artifacts but computational ones – they are **execution traces** that anyone (or any machine) can replay to be convinced. This foundational stance will guide us through the rest of this chapter and book, as we reconstruct the edifice of mathematics on top of a binary, verifiable substrate.

1.2 Deconstructing Mathematical Primitives: A Binary Foundation

Having outlined the philosophical stance, we now *deconstruct the basic primitives of mathematics* (number, set, logic, function, etc.) and reinterpret them through the binary-executable lens. The way these concepts are traditionally introduced often reflects historical development or pedagogical convenience, prioritizing human intuition over computational rigor. Machine Epistemology necessitates rebuilding these concepts from the ground up, starting with the binary substrate and insisting on explicit executability and verifiability.

A. The Nature of Number: Beyond Counting

Traditional Intuition: The natural numbers (0, 1, 2, 3, ...) are usually introduced as self-evident entities for counting discrete objects or ordering. Peano's axioms formalize them with the existence of a base element

0 and a successor function $S(n) = n+1$, and properties like induction. Set-theoretic foundations even define numbers in terms of nested sets (Von Neumann ordinals: $0 = \emptyset$, $1 = \{0\}$, $2 = \{0,1\}$, etc.). These approaches take “**number**” as an initial given concept or build it from an abstract notion of set, largely *ignoring the computational substrate*. Induction is presented as a logical principle rather than a physical process. In practice, however, when we use numbers in computing, we immediately represent them in binary (or some base) and implement operations algorithmically. This gap between formal axioms and actual computation is where we apply our new perspective.

Binary Foundation View: We reject the notion of numbers as inherently primitive objects – instead, the only true primitive is the ability to distinguish two states (the bit). A **number** is thus understood as a *constructed binary pattern together with algorithms* that operate on such patterns. For example, what is the number 1? In a binary computer, it could be an 8-bit pattern `00000001` in memory, meaning it’s the first non-zero state after the zero pattern `00000000`. The number 2 is then `00000010`, etc. In general, an n -bit binary string can represent numbers from 0 up to $2^n - 1$. In this view, 0 is not an abstract “nothing” but corresponds to a *zeroed memory state* – a specific configuration of bits all set to 0. The number 1 corresponds to the simplest *distinguishable signal* from 0 (e.g., a single 1 in the least significant bit). The successor operation $S(n)$ (which yields $n+1$) is no longer an abstract axiom but a concrete procedure: e.g., a **BDI** node performing an increment on a binary register, flipping bits and handling carries. This is an algorithm encoded in the Binary Decomposition Interface graph or in machine code – a finite state transformation that we can execute. Addition and multiplication are likewise specific circuits or programs (which can be represented as compositions of BDI operations like `ARITH_ADD`, `ARITH_MUL`, etc.). In this way, **Peano’s axioms become descriptive** of how our constructed number system behaves, rather than fundamental truths. The axiom “every number has a successor” corresponds to the fact that our algorithms can always produce a next binary state (until physical memory limits). The principle of mathematical induction – usually a schema for proofs – is here interpreted as a property of *iterative computation*: if a certain property holds for 0 and you have a verified procedure to go from n to $n+1$, then by running that procedure repeatedly, it will hold for all representable n . Induction is guaranteed by the reliability of the hardware or the BDI execution model in applying the successor operation repeatedly, and by the absence of unexpected halting. In short, we **build** the natural numbers from the bit up, rather than assuming numbers as an unexplained given. This grounding in the binary makes numerical truths a matter of verifiable algorithmic processes (for instance, a statement like $2+3=5$ is true because if you load the binary for 2 and 3 into an adder circuit or BDI graph, the output is the binary for 5, which can be checked).

This computational re-foundation of number aligns with the concept of *Kolmogorov complexity* in algorithmic information theory, which essentially defines the complexity of an object as the length of the shortest binary program that produces it. A number like 1 or 2 has low Kolmogorov complexity because there are very short programs (or descriptions) that produce it (indeed, we can produce any fixed natural number with a short addition loop or bit pattern). The philosophical point is that **a number is only “known” to the extent that we have an explicit method to represent and generate it**. In classical math, one might say “there exists a number with such-and-such property” without constructive content. Machine Epistemology pushes us toward a constructive stance: existence is tied to constructibility within the binary framework. If someone claims a very large number has a certain property, our framework asks for a program (or BDI graph) that can *verify* that property for that specific number, even if by exhaustive or probabilistic methods. This does not mean we abandon classical mathematics, but we reinterpret its claims in a way that they are subject to verification by explicit computation.

In practical terms, when implementing mathematics in a computer, this view is already standard. What we add is the insistence that **this is not merely implementation, it is ontology and epistemology**: the *reality* of the number 10, for instance, is nothing more (and nothing less) than the array of bits 1010_2 along with the operations that can be performed on it (increment, add, multiply, etc.), all of which are ultimately reducible to switching circuits or BDI node operations. This demystifies numbers and prepares them for integration into an **executable mathematics** framework.

B. The Nature of Sets: Beyond Abstract Containers

Traditional Intuition: Set theory is often treated as the foundational language of mathematics (Zermelo–Fraenkel set theory with Choice – ZFC – being the de facto foundation for classical mathematics). In naive terms, a set is described as a "collection of distinct elements," an abstract container that can hold any objects. We write $x \in S$ to denote membership. Operations like union $S \cup T$, intersection $S \cap T$, complement, power set, etc., are defined logically (e.g. $x \in S \cup T$ iff $x \in S$ or $x \in T$, etc.). While this abstract notion is powerful, it also led to well-known paradoxes (Russell's paradox about "the set of all sets that do not contain themselves") that forced a careful axiomatization. Importantly, the usual conception of an infinite set or an arbitrary subset of an infinite set is *highly non-constructive* – we often quantify over uncountably many possibilities with no algorithm to list them. This is acceptable in classical math, but from a computational standpoint it's problematic since a machine cannot even finitely represent an arbitrary infinite set or perform an unbounded search to decide membership if the set is given implicitly.

Binary Foundation View: In our framework, a "set" must be something with a concrete **representation in memory and executable procedures for its operations**. In other words, we treat a set *like a data structure*. Within the BDI system (and a prototype system called Chimera built on BDI), a set is not an amorphous platonic bag but, for example, a **tagged BDI memory region** designated to store certain elements. The membership relation $x \in S$ is realized by an *algorithm* that searches this region to check if the binary representation of x is present. There are multiple possible implementations: sets can be implemented as hash tables, where membership is a hash lookup; or as sorted binary trees, where membership is a tree search; or as bitvectors, where membership is a single bit index check. Each of these is a well-understood data structure in computer science, and each can be composed out of BDI operations (for example, a tree can be made of node structures with pointers implemented as references to memory addresses, comparisons implemented by arithmetic or logic BDI nodes, etc.). **Union** of two sets S, T becomes a specific algorithm that merges two data structures (if hash tables, it might iterate over one and insert into the other; if bitvectors, it could be a bitwise OR operation on two binary strings). **Intersection** is similarly an algorithm (hash table: iterate and check, bitvector: bitwise AND). Even an infinite set, in order to exist in our system, must be represented by some rule or generator that can enumerate its elements to whatever extent required. For instance, the set of all prime numbers could be represented by a program (or an infinite stream generator in BDI) that can produce primes on demand. Without such an operational content, a set is not considered fully defined.

This view aligns with the **constructive** or **computable set theory** viewpoint in logic, where one distinguishes between a set given by a membership algorithm versus an arbitrary non-computable set. In classical ZFC, one can assert the existence of non-computable sets (indeed, most real numbers form an uncountable set, almost all of which are non-computable). Machine Epistemology would be skeptical of considering such objects as "knowledge," since no machine can ever verify properties about them in general. Instead, we focus on sets that are *recursively enumerable* or decidable or otherwise tied to computational mechanisms.

By reifying sets as data structures, we also gain the ability to attach **invariants** and **proofs** to them as part of their metadata. For example, a set in BDI might carry a **characteristic function** – a BDI subgraph that takes an input x and returns a boolean (0/1) indicating membership. This is effectively the set's definition as an executable function. A rich BDI metadata system could allow us to state that a certain BDI subgraph is the membership test for a certain memory region, and even formally verify that property. This moves us closer to the ideal of **proof-carrying data structures**, where, e.g., a balanced tree might carry a certificate that it is balanced (perhaps a recursive proof attached to each node) and the operations on it carry proofs that they preserve the balance. Those ideas resonate with research in formal verification and certified programming (such as projects in Coq/Agda where one writes programs with proofs of correctness that the compiler checks). For instance, one could cite the work on fully verified red-black trees or union-find structures in proof assistants – these ensure that every operation on the data structure maintains certain set-theoretic specifications.

In a BDI-based mathematics, **the set concept becomes concrete and finitary**. Any set we work with is either finite (in which case an explicit memory representation exists) or it's given by an intension (like a formula or generator) which is itself executable. This does not prevent us from discussing infinity; it only means whenever we claim something about “all n in \mathbb{N} ” we acknowledge this as shorthand for “for each n , our algorithm/proof works for that n in a uniform way.” The generality is captured by a program schema rather than an actual completed infinity. As a result, paradoxes like Russell's are avoided not by complex axioms, but simply by the fact that we cannot instantiate an impossible self-referential set in actual memory. Any attempt to construct “the set of all sets not containing themselves” in BDI would either not type-check or not terminate, so it wouldn't be a valid constructible set.

This grounded view of sets underscores a theme: **mathematical existence is tied to constructability**. A set exists for us if we can build it or at least simulate it within the BDI substrate. This idea mirrors the philosophy of **constructivism** in mathematics (espoused by Brouwer, Bishop, etc.), but here it's less about philosophical preference and more about *practical enforceability*: a non-constructive existence claim simply cannot be verified by any machine, so in Machine Epistemology it has no epistemic weight until made constructive.

C. The Nature of Logic and Proof: Beyond Symbolic Derivation

Traditional Intuition: Logic is the framework that underpins mathematical reasoning. In classical propositional and predicate logic, we have formal systems of inference (e.g. truth tables for connectives, natural deduction or sequent calculus for proofs). A statement is *logically valid* if it is true under all interpretations (tautology), and a sequent (collection of premises leading to a conclusion) is valid if there is a formal proof. Traditionally, logic is taught as an abstract symbol game: one manipulates formulas by applying rules like modus ponens, introduction/elimination rules for \forall , \exists , \wedge , \vee , \rightarrow , etc. The emphasis is often on **consistency and provability** within an axiomatic system. A proof is a static object – a sequence of formulas each justified by a rule – and its validity is a matter of each step following from previous ones.

Binary Foundation View: We reconceptualize logic itself as an *executable system of state transformations*. In essence, **logic is treated as just another DSL**, one whose domain is statements about states, and whose inference rules are akin to operations. A logical proposition can be seen as a specification of a desired output of a computation (true or false), and a **proof is literally a program** (as touched on earlier with Curry-Howard). Within BDI, we could have a library of logic operations: for example, a node type for

applying modus ponens, or a node type that checks a truth table of a formula. However, more directly, we note that any finite logical inference can be directly **executed by brute force** at the bit level. For instance, checking a propositional tautology P means evaluating P under all 2^n possible assignments of its n atomic propositions – this is a finite (if potentially exponential) computation. Similarly, first-order logic statements can be checked up to a finite model size or through search (e.g. SAT and SMT solvers are essentially *programs that find proofs or countermodels* for logical formulas). The success of SAT solvers and automated theorem provers in recent decades reinforces the idea that **proof search can be mechanized** and run on silicon. In our framework, a proof is not a static derivation but an *execution trace* that can be replayed and verified step by step. We imagine encoding proofs as graphs: each inference rule application is a node, its premises are incoming edges, and its conclusion is an outgoing edge, forming a directed acyclic graph that ends in the theorem to be proven. Such a proof graph is essentially a program that, when “run,” yields the theorem as output from axioms as input.

By bringing logic down to binary operations, we also integrate **verification as part of logic**. A proof in BDI can carry a cryptographic *ProofTag* – essentially a hash of an externally verified logical derivation. A special BDI operation `VERIFY_PROOF` can then check that the execution trace corresponds to that hash. This is akin to a *proof-carrying code* approach for mathematics: the proof comes with a certificate (the hash or some concise evidence) that it matches a known correct derivation (for example, one produced by Coq or Lean), and the machine can verify the certificate. This addresses a potential worry: how do we trust that the BDI’s proof procedures are sound? The answer is we can cross-link them with traditional formal methods. For example, each BDI inference node for group theory or calculus can reference a lemma in a proof assistant, ensuring that the step is semantically correct. The BDI system could optionally check these references (maybe not every time, but in audit mode) to guard against bugs in its own inference implementation.

An important consequence of our view is that the **distinction between “proof” and “program” evaporates**. Proofs *are* programs, and proving a theorem is equivalent to exhibiting a program that produces a certain output (truth) from certain inputs (axioms). This means that techniques from software engineering become applicable to mathematics and vice versa. For instance, one can debug a proof by stepping through it, or optimize a proof by finding a shorter sequence of steps (similar to optimizing code). The notion of proof complexity becomes akin to time complexity of an algorithm: a proof of length N is like an $O(N)$ algorithm to verify the theorem. This is illuminating in contexts like cryptography or complexity theory – e.g., NP versus P can be seen as asking whether every succinct proof (NP certificate) has a short executable proof verification (P algorithm). In our epistemology, a theorem’s difficulty is measured by the computational resources needed to verify it.

By grounding logic in executability, we also clarify the concept of **consistency**. Normally, consistency of a theory means you cannot derive a contradiction. In our terms, a theory (a set of axioms and inference rules) is consistent if there is no sequence of executable steps that yields a false statement from no premises. This could in principle be checked by a model search or by attempting to derive an inconsistency (like a SAT solver looking for a satisfying assignment to the negation of a supposed theorem). While general consistency is uncheckable (by Gödel’s second theorem, a sufficiently strong theory cannot prove its own consistency), our emphasis is less on absolute consistency and more on **local verifiability**: each proof stands on its own as an executable artifact that either works (verifies) or doesn’t. The usual Gödelian issues do not negate our approach; rather, they indicate there will always be true statements our system cannot prove with given resources, but those statements are not “known” in our sense until a proof (program) is found.

In summary, logic in the Machine Epistemology view is the science of reliable transformations on binary-encoded knowledge. It provides the rules (operations) by which we can validly transform one set of bit patterns (premises) into another bit pattern (conclusion) while preserving semantic truth. We demand that each such transformation be not only *sound* (truth-preserving) but also *executable* and *traceable*. This ensures that whenever we assert a logical result, we have a concrete runtime evidence for it. The philosophical payoff is significant: it eliminates the need for “**mathematical intuition**” as a mysterious faculty – instead, what we call intuition can be seen as an internalized sense of what computations will succeed or fail. By making everything explicit, we allow machines to partake in what was once exclusively a human endeavor, and indeed to outperform humans in many aspects (as modern theorem provers already can beat humans in propositional logic or even in certain algebraic proofs).

D. The Nature of Functions and Computation: Beyond Black-Box Mappings

Traditional Intuition: Functions are fundamental objects in mathematics – a function $f: X \rightarrow Y$ is typically defined as a rule that assigns to each $x \in X$ an element $f(x) \in Y$. In set theory, a function can be seen as a special set of ordered pairs $\{(x, y) : y = f(x)\}$. Often, functions are treated as black boxes: we care about their input-output behavior (extensionality) but not necessarily the internals. For example, the function $f(n) = 2n$ (doubling) is the set of pairs $\{(0,0), (1,2), (2,4), \dots\}$ or given by the formula $2n$. In computation, however, a function is an **algorithm** or a circuit – there is a clear notion of how input data is transformed into output data. The lambda calculus and Turing machines formalized the idea of a function as something computable, but pure mathematics often speaks of functions that are not computable or not explicitly given (e.g. “there exists a function with such property,” without constructing it).

Binary Foundation View: We assert that every meaningful function in our system must have an **executable description**. A function is essentially a DSL element that corresponds to a subgraph of the BDI graph – a *callable component* with input ports and output ports defined. Indeed, BDI supports an operation type like `CALL` or function application as a node, which can invoke one subgraph from another. Thus, a function $f: X \rightarrow Y$ would be implemented as a BDI subgraph (with perhaps a label or identifier) such that providing the binary encoding of an $x \in X$ at its input leads, after execution of the subgraph, to the binary encoding of $f(x)$ at its output. The **extensionality** of a function (the set of input-output pairs) is then a *derived property* of this executable: one could in principle run through all possible inputs (if finite or bounded) to collect the pairs. But the primary definition is intensional – the algorithm itself. In this sense, *function* and *algorithm* are synonymous in our framework, aligning with the Church-Turing conception of a computable function.

This has several implications. First, uncomputable functions (like a random oracle or a truly arbitrary mapping without pattern) have no place unless they are approximated by a procedure. For example, in classical analysis one often discusses functions like $f(x) = 1$ if a certain axiom is independent and 0 otherwise – such a function is undefinable because it encodes a non-computable predicate. We would simply say that “function” is not a valid object of knowledge, since no machine can realize it. On the other hand, any function that we define with a simple formula or case distinction *is* computable in principle (though possibly expensive). For instance, the Riemann zeta function $\zeta(s)$ can be defined by a series – we know how to approximate it to any given precision, so it is algorithmic knowledge. But a function defined by an existence theorem might not be.

Second, functions in BDI carry **metadata for verification**. If we claim $g(x) = h(x)$ for all x (two functions are equal), we are expected to provide either a proof (in the form of a BDI subgraph that, given any x ,

shows the outputs of g and h are identical) or enforce it by construction (maybe g and h are actually the same subgraph referenced by two names). The system could allow *ProofTagging* of functional properties similar to how it handles proofs. For example, a `ASSERT` operation in BDI (as in `META_ASSERT`) might assert that two outputs are equal, and if this assertion is annotated with a proof, the BDI verifier can check it.

Finally, this view ties back to the idea of **structured binary decomposition** mentioned earlier: any complex function or system can be decomposed into simpler binary operations. This is exactly what programming (especially in a low-level language) does. Our claim is that mathematics is *no different* – say you have a group homomorphism $f: G \rightarrow H$, you might traditionally only specify it by a few properties ($f(xy) = f(x)f(y)$). But to truly *know* this homomorphism, one ultimately writes algorithms (or at least explicit formulae) for computing f on generators of G . The explicitness can be deferred by abstraction, but it cannot be eliminated if we demand verifiability. If someone posits a function defined nonconstructively (like choosing an element from each equivalence class of an infinite relation – the Axiom of Choice scenario), our framework would treat that as an open problem: "please provide an algorithm for making those choices, or you haven't given me knowledge, only a conjecture that some knowledge exists." This again echoes constructive mathematics' stance but is enforced here by the requirement of binary executability.

In summary, we deconstruct the notion of function from an arbitrary mapping to a *process*. Every function is an executable *transformation of binary states*. Composition of functions corresponds to connecting the output of one subgraph to the input of another. Properties like injectivity, surjectivity, continuity (if we talk about topological notions) etc., all must be translated into conditions that can be checked or at least pursued by computational means. For instance, continuity in a computable analysis sense means: there is an algorithm that, given an input and an ϵ , finds a δ such that the condition holds. The framework thereby pushes us toward constructive definitions of all such properties.

By rebuilding number, set, logic, and function on explicit binary operations, we set the stage for a mathematics that is **automatically ready for machine use**. This does not trivialize mathematics; rather, it forces precision and eliminates ambiguity. Each concept carries both a *meaning* and a *method*. The meaning might align with classical definitions in the limit of infinite resources, but the method ensures that whatever we assert, we could in principle demonstrate with a finite computation. This tightly couples the *semantics* of mathematical statements with their *computational content*, which is a theme that will recur throughout our development.

1.3 Mathematics as a Hierarchy of Domain-Specific Languages (DSLs)

One of the core insights of this framework is that all the diverse branches and systems of mathematics can be viewed as a hierarchy of **Domain-Specific Languages (DSLs)**, layered over the binary substrate. In computer science, a DSL is typically defined as "a programming language or executable specification language that is tailored to a specific application domain" ⁵. Examples include HTML as a DSL for webpage layout, SQL for database queries, or Verilog for hardware circuits. Here we broaden the notion: **each mathematical theory or system (arithmetic, linear algebra, group theory, topology, etc.) is considered a DSL**, complete with its own syntax (notations and symbols), semantics (meaning of those symbols in terms of operations or in a model), and pragmatics (how humans use it to solve problems). Under this view, *ZFC set theory* is a DSL in which one can encode virtually all of standard math; Peano

arithmetic is a DSL for natural numbers; Euclidean geometry can be seen as a DSL with its own axioms and constructions.

Why is this a useful perspective? First, it emphasizes that *mathematical languages are engineered artifacts*. Just as a programming language is designed with certain goals (readability, efficiency, safety) in mind, mathematical formalisms were developed by humans to tackle certain classes of problems with ease. For example, algebraic notation was developed to succinctly solve polynomial equations, and it's essentially a DSL atop basic logic. Category theory can be seen as a DSL that abstracts patterns in mathematical structures. By seeing them as DSLs, we recognize that we can **translate** or **compile** between them when needed – much like a compiler translates a high-level DSL into machine code. Indeed, a significant part of advanced mathematics is finding translations between different frameworks (e.g. using analytic methods to solve a number theory problem, which is like using the DSL of calculus on a problem stated in the DSL of integers).

In our Machine Epistemology framework, the DSL view is literal: each mathematical theory *will be implemented as a module on top of BDI*. We imagine a stack where at the bottom is the BDI's native binary executable environment; above that, there might be a core language for logic and set theory; above that, one for numbers and algebra; above that, maybe one for analysis, etc. This stack is not strictly linear – it could be more of a web of DSLs that interoperate. But crucially, the **semantics of every DSL is given by a mapping into the BDI layer**. This is analogous to how the semantics of SQL can be given by its translation into lower-level file system operations, or how the semantics of a high-level language is ultimately given by what sequences of CPU instructions (binary) it corresponds to. In the case of mathematics, providing semantics means providing an interpretation of the mathematical statements as specific configurations and executions in the binary substrate. For example, the statement “for all $x \in \mathbb{N}$, $x + 0 = x$ ” in a DSL for arithmetic would map to a BDI graph that takes an input encoding x and, via an addition circuit with 0 , produces an output identical to the input – and then a verification that those bits are equal for all possible x (which might be shown by induction as another BDI sub-program).

The DSL perspective also sheds light on **symbolic compression**, which was mentioned in the abstract as “structured symbolic compression over binary distinctions.” Each DSL serves as a *compression mechanism* for certain patterns of binary operations. For instance, instead of reasoning at the bit level about toggling transistors to add two numbers, we introduce the symbol “ $+$ ” and laws like commutativity to reason abstractly – this is a huge compression of information, since one symbol $+$ encapsulates an entire class of circuits or algorithms. Similarly, the notion of a group in algebra compresses a lot of information (a set with an operation satisfying certain axioms) into a single concept, allowing us to reason at a higher level and ignore the lower-level details momentarily. However, unlike traditional abstract mathematics, we never lose sight that this is compression, not mystification: we always have the ability to decompress the symbols back into the binary operations when needed. This is analogous to how modern software engineering might use high-level libraries but can inspect the assembly code if debugging at the lowest level is needed.

Because mathematics is a tower of DSLs, **metalogical and metamathematical questions** become analogous to compiler theory and interoperability of languages. For example, consistency of a theory is like saying the DSL has a sound semantics in the binary world (no program compiled from it will produce a contradiction state). Relative consistency or relative computability between theories would correspond to whether one DSL can be translated into another. Famous results like Gödel's incompleteness can be seen through this lens: any sufficiently powerful DSL (capable of arithmetic on its own formulas) cannot prove all truths expressible in it – which is somewhat analogous to the halting problem in computation (no single

programming language can decide all properties of programs written in it). This analogy is more than superficial: Gödel's arithmetization of syntax essentially *reduced mathematics to computations on binary strings* (the Gödel number encoding), which is precisely our starting point. In a sense, Machine Epistemology is "Gödel-complete": we accept that mathematics, being equated with computation, inherits both the power and the limitations of computation. But this is not a defect; it's a reality to work with. For instance, there will be true statements our system can't automatically prove, but we might recognize them via meta-arguments and then extend the system (like adding a new axiom or oracular operation, akin to upgrading a programming environment with a new primitive).

By framing mathematics as DSLs, we can apply **formal language theory** and **compiler techniques** to mathematics. A practical benefit is that we can use existing tools to implement these DSLs. Indeed, projects in automated theorem proving often create custom languages (e.g., SMTLIB language for SMT solvers, or idioms within proof assistants) to express problems. Our approach would unify these with the underlying BDI, meaning the distance from abstract specification to execution is greatly reduced. Traditional pipelines go: *math problem* \rightarrow *formalize in high-level logic* \rightarrow *manually or automatically derive proof* \rightarrow *trust the proof or check it with a proof checker*. We propose: *math problem* \rightarrow *represent directly as a BDI executable graph (a DSL instance)* \rightarrow *run it or model-check it on BDI to verify*. The BDI acts as a universal *semantic execution engine* for all these DSLs. In a way, BDI is like a "mathematical virtual machine" that understands the semantics of high-level mathematical operations because those semantics are given to it in a structured way (through metadata and operation mappings).

To illustrate, consider a hierarchy: At the lowest level DSL0 we have **Boolean logic and bit operations** (the native language of BDI). Next, DSL1 might be **Peano arithmetic** (PA) – its syntax is numbers and $+$, $*$, \times , etc., and its semantics is given by how those operations compile to bit operations. Above that, DSL2 could be **field theory** which includes rational numbers, etc., which compile down to arithmetic on integers and hence to bits. Another branch, DSL1' could be **set theory** – whose semantics might be given via a set of memory regions in BDI as discussed. Yet another, DSL2' could be **first-order logic** – whose semantics is essentially truth evaluations, which at the base require iterating over cases or using algorithms. As we go higher, DSL3 could be a combination like **elementary analysis** (real numbers and calculus) which might be built on top of rational approximations plus some limit procedures. Each step up introduces new abstractions which are *verified by the layer below*. If something goes wrong (say we allow a paradoxical combinatorial object), it would manifest as an inability to ground it in the lower level (like a program that doesn't compile).

This layered approach is similar to how a complex software system has multiple layers (machine code, assembly, C language, libraries, scripts, etc.), each verified on top of the other. In our context, because each layer is still within a unified formalism (all layers ultimately end up as BDI graphs), cross-layer reasoning is possible in a way not usually done in mathematics. For instance, one could formally verify that a certain group theory proof (done in a high-level algebraic DSL) indeed corresponds to a certain low-level logical derivation by following the chain of compilation. This could lead to a new kind of **meta-proof** or *proof traceability*, where a proof in a high-level domain carries along a certificate that it is sound by referencing lower-level proofs or operations. BDI's metadata handle is designed to support this kind of rich linkage.

In conclusion, seeing mathematics as a collection of DSLs underscores our theme that **mathematical knowledge is engineered and compositional**. It is built up in layers of abstraction, much like software. Each layer provides a compression of complexity that makes it easier for human or AI mathematicians to work, but none of these layers float freely – they are all grounded in the binary executable reality. This

perspective will guide how we design the Binary Decomposition Interface to support multiple semantic layers, and how we ensure that when mathematicians (human or machine) work at a high level, they do so with the guarantee that a low-level, checkable realization exists for every step they take.

1.4 Symbolic Compression, Information, and Verifiability

A critical advantage of the Machine Epistemology framework is that it provides a natural marriage between **mathematics and information theory**. We have mentioned “symbolic compression” as the role of DSLs: compressing complex low-level sequences into higher-level symbols. Here we delve deeper into the idea that mathematics itself can be seen as a process of finding **compressed representations of patterns** in information, and that this compression is only valuable if it is lossless with respect to verifiability (i.e., one can always decompress to check correctness). This aligns with the principle of **minimum description length** in learning and the notion of **Kolmogorov complexity**, which we touched on.

In information-theoretic terms, a theorem or a formula can be thought of as carrying information (Shannon entropy, perhaps) about the models or numbers it pertains to. A proof, then, is a method to reduce uncertainty about the truth of that formula by providing a *procedure* to verify it. The shorter (or more efficient) the proof, the more it compresses the reasoning needed. In an ideal sense, a perfect proof is one that compresses an exponential search (trying many possibilities) into a polynomial sequence of steps – akin to finding a neat algorithm for a problem instead of brute force. This connects to the concept of **computation as the engine of knowledge**: if you can compute something in fewer steps, you have in effect increased your knowledge efficiency.

Our binary framework allows us to quantify such things. For a given statement S , consider the space of all possible bit strings of length N as potential proofs (or disproofs). A priori, if we had no guidance, finding a proof is like searching in that space – expensive. But mathematical structure (lemmas, heuristics, intuition) guide us to a much smaller subset of candidates. Those guides are DSL abstractions that significantly reduce the entropy of the search space. In Machine Epistemology, we can attempt to measure the **entropy of a mathematical problem** and the **entropy reduction achieved by a proof**. This is a novel angle: a proof not only convinces, it *compresses* the verification task. If a proof is very convoluted and long, it hasn't compressed much (perhaps the problem is inherently complex). If a proof is elegant and short, it's a high compression of a lot of brute-force checking.

Kolmogorov complexity tells us that the absolute information content of a theorem is essentially the length of the shortest program that outputs a proof of it (or verifies it). If our framework is realized, that “shortest program” can be taken quite literally as a BDI executable. For example, the assertion that a certain large number is prime has high intuitive complexity if approached by trial division, but the Pratt primality certificate provides a short proof (relative to the number's size) that can be verified quickly – it compresses the work. Similarly, the four-color theorem's computer-assisted proof can be seen as a huge compression of checking many cases; it was done by a program rather than enumerating by hand.

This leads to an interesting identification: **Mathematical insight = finding the right compression of a problem**. In our terms, that means finding the right representation in a higher-level DSL such that the problem's structure becomes clear and the solution steps are fewer (i.e., the proof is shorter). For instance, using group theory (a DSL) to solve a combinatorial puzzle compresses the reasoning by leveraging symmetry. This viewpoint dovetails with how mathematicians informally talk about deep ideas making proofs simpler or more conceptual – those ideas are compressions.

However, unlike lossy compression, here we require lossless compression: nothing of truth value is lost. So any time we compress (abstract), we ensure we can decompress (instantiate) if needed. This is taken care of by the DSL semantics mapping to BDI. If someone doubts a high-level proof, we can always machine-check it down to binary. Think of it as a zip file containing all the low-level details, with the high-level proof being like a summary. BDI can unzip it and check every bit.

This interplay also has a thermodynamic flavor. When we say a true statement corresponds to a *low-entropy state*, we imply that the system has resolved the uncertainties and arrived at a stable configuration (the verified proof graph). During the process of proving, the "entropy" (uncertainty, number of possibilities) goes down. If one wanted to push the physics analogy, an elegant proof is like a reversible computation (little entropy produced), whereas a brute-force proof is like a lot of heat dissipated.

Now, consider **verifiability**. In information terms, verification is a channel by which the information of truth is transmitted from the proof to the verifier. A proof's reliability can be enhanced by redundancy (like error-correcting codes) – e.g. having multiple independent proofs, or a proof that cross-checks itself. In BDI we mentioned the idea of ledgering: hashing each step and possibly recording it on an immutable ledger. This is adding a layer of security: even if someone tried to tamper with a proof, the hashes wouldn't match, revealing the tampering. It's analogous to how digital communication adds checksums to ensure integrity of data. So our system naturally can integrate cryptographic techniques to ensure that once a proof is verified, it stays verified (nobody can alter a bit without being noticed). This is crucial if mathematics becomes a machine-driven collaborative enterprise; we need trust that results aren't corrupted. The use of cryptographic **ProofTags** means each proven theorem can carry a signature that any future user can check, effectively certifying the theorem within this ecosystem.

We also consider **entropy in learning and AI** – although this chapter is about mathematics, one motivation is intelligent systems. An intelligent system, from our perspective, is one that can compress and generalize patterns (learn) and then act on them. By establishing mathematics as executable knowledge, we pave the way for machines not just to compute but to genuinely *learn mathematics or other structured knowledge* and verify what they learn. The principles we lay (binary grounding, DSL hierarchy, proof as execution) are equally the principles needed for an AI to explain and justify its decisions. If an AI derives a conclusion, we want an **executable trace** of how, not just a mysterious neural network output. BDI and Machine Epistemology push for that transparency.

In formal verification terms, our approach is aligned with the concept of **proof-carrying code** and **certifying algorithms** – algorithms that not only produce an output but also a certificate that the output is correct. Many algorithms in computer science have been enhanced this way (for example, a spanning tree algorithm that outputs the spanning tree and a proof that it's of minimum weight). We essentially treat mathematical problem solving as a certifying algorithm: solving a problem is inseparable from producing a certificate (proof). If it doesn't produce a verifiable certificate, it's not a valid solution in our book. This would revolutionize mathematical publication – results would come with machine-checkable proofs by requirement (something the QED manifesto advocated ²).

To summarize this section, mathematics in our framework is about **finding maximally efficient (compressed) binary representations of truths**. The success of a mathematical theory can be partly measured by how well it compresses a class of problems (for instance, arithmetic compresses repeated addition into multiplication, calculus compresses infinite processes into finite rules via limits, etc.). But we always guard this compression with the ability to decompress (verify). In doing so, we tie together the

notions of **simplicity, informativeness, and certainty**. A good proof is simple (short description), informative (high explanatory power), and certain (fully checkable). These are not just subjective virtues but quantifiable features in an information-centric view.

In a way, we are fulfilling some of Leibniz's dreams: he imagined a *characteristica universalis* and *calculus ratiocinator* – a universal logical language and a method to compute truth values, so that disputants could just say "Calcuemus!" (let's calculate) to resolve a question. Our Binary Mathematics is an attempt at that: a universal executable semantic language for all knowledge. With it, we can indeed say "let's compute and verify" when confronted with a conjecture. The "calculation" may be very heavy (requiring supercomputers or clever heuristics), but it is in principle a mechanical task, not requiring mystical insight once set up. This puts mathematics and science on the same footing as any engineering discipline – we build verifiable artifacts (proofs and programs) to demonstrate results. What differentiates mathematics is just the level of abstraction and generality of those artifacts, but not the nature of verification.

1.5 Truth as an Executable, Verifiable Outcome

Having laid the groundwork, we sharpen the concept of **truth** in mathematics under Machine Epistemology. Traditionally, truth of a mathematical statement is understood in model-theoretic terms (a statement is true in a structure if it holds under the interpretation) or proof-theoretic terms (true if provable from axioms). We offer an alternative that effectively merges these: truth is *that which a correctly constructed computation can produce as output*. More precisely, given a statement S in a DSL, we say " S is true" if there exists a *verification program* for S that, when executed on an appropriate computing device (our BDI machine), halts and produces an affirmative result, and no counterexample can be produced by any similar program. This makes truth inherently **relative to the formal system (DSL) and the substrate**. It's not an absolute notion floating in metaphysical space, but a conditional notion: true *in* system X , under encoding Y , via substrate Z .

Concretely, consider a statement like " $\forall x \in \mathbb{N}, P(x)$ ". In our framework, to say this is true means we have a BDI program that takes an arbitrary natural number n (in binary) and returns `true` (1) as output, and we can formally show (or the program itself demonstrates by induction or other means) that it will always return true for any n . This could be an inductive proof encoded in BDI: a base case check for $P(0)$, and a loop or recursive call that shows $P(n)$ to $P(n+1)$. If such a program exists and can be verified (perhaps by model checking or by a meta-proof attached), then the statement is considered true. The "for all" has been turned into a universal verifier that one could run on any given n to check $P(n)$, plus a proof that you don't need to run it for all n because the result is generic.

Truth thus becomes **operational**: it is linked to the success of an algorithm. If a statement is undecidable (like the continuum hypothesis in ZF set theory), it means no such BDI program exists within the system (either one that verifies it or one that refutes it). It remains neither true nor false in that context until the system is extended (like adding CH as an axiom would permit a trivial program that just returns true for it).

One might wonder: does this just collapse to the proof-theoretic notion of truth (provability)? In a sense yes, but we broaden what counts as a proof. It need not be a proof in a human-oriented logic system; it can be any computation whose outcome entails the result. For example, to show a large even number is prime or composite, we don't require a human-style proof; a certified computational test (like a Pratt certificate for primality, or an actual factor found for compositeness) suffices as truth evidence. So truth in our sense is closer to **"there is a verification procedure that confirms this"**. This even covers statements about

physical reality insofar as they can be verified by experiments (though here we focus on mathematical truth).

In philosophical terms, this position is a form of **verificationism** but via machine. It echoes the pragmatist and intuitionist idea that truth is what we can in principle verify. However, unlike naive verificationism, we allow that truths might exist that are just too complex to verify with current resources – but then those truths are simply unknown, not yet part of the body of mathematical knowledge, which aligns with common usage (an unproved conjecture is not part of “known truths” of math). If one day a machine finds a proof, it becomes a known truth.

By making truth executable, we also ensure **consistency** is built-in: a false statement is one for which a counterexample can be computed. For instance, “there exists x such that $P(x)$ ” is false if a program can verify $\neg P(x)$ for all x (or find that for each x , $\neg P(x)$ holds). If a system erroneously proved both S and $\neg S$, then we would have programs for both a result and its negation, which when executed together reach an obvious contradiction (like outputting 1 and 0 for the same query). In a well-designed BDI environment, this should be detectable – perhaps by a consistency-checking routine that flags when two proof graphs prove opposite outcomes, akin to finding an unsatisfiable core. In practice, our approach should minimize such risks by construction, since BDI operations come from sound rules and algorithms.

One can ask: what about statements that aren’t about concrete computations, like “There are infinitely many primes”? How is that executable? In our framework, “infinitely many primes” is an abbreviation of a family of statements: for each n , “there is a prime larger than n ”. To verify “infinitely many primes”, we would likely provide an algorithm that given any n , can produce a prime larger than n (Euclid’s proof is essentially that algorithm: take primes up to n , multiply them plus 1, find a prime factor). Indeed Euclid’s proof can be seen as a function $f(n)$ that outputs a prime $> n$. Formalizing that in BDI would involve the construction of that output and a verification it’s prime. So even existence of infinite sets or unbounded truths is handled by an algorithmic schema.

Thus, “true” means “**we can always get a yes answer by computation, and never get a counterexample by any computation**”. This makes truth a subset of what’s classically true (since classical allows truth that we might never know). We are essentially adopting the stance that if something is true but inherently unverifiable in any computational way, it might as well not concern us – it’s metaphysical fluff. All that matters for science (and certainly for applications like AI) is the truths we can establish or use effectively.

It’s worth noting that this approach relates to **formal verification in hardware/software**. When an engineer says a chip design is correct, they mean they ran a model-checker or theorem prover which exhaustively or symbolically verified it against a spec. That’s “true by verification”. If the spec had some property that was true but the tools couldn’t verify it, the engineer wouldn’t rely on it. Similarly, our mathematics will be a **specification language** with the guarantee that anything proven in it has been model-checked down to binary logic. This could be seen as an extreme realization of the QED vision ² : not only are proofs checked, but the notion of truth is *defined* as “has been checked (or is checkable)”.

One might worry this limits mathematics – do we allow only things that machines can handle? In practice, the answer is that machines (with human guidance) can handle a vast and growing portion of mathematics. Even very abstract fields like category theory or higher-dimensional topology can be (and have been) encoded in proof assistants. There will be a creative aspect in finding the right encodings (much like writing good software), but none of the content is lost. We can still discuss inaccessible cardinals or exotic

geometries, but as DSLs that require perhaps some augmentation of our computing power (like oracles or higher-type computation). As long as we define the semantics clearly (even if it's "assume an oracle for this"), it's within our framework.

In summary, we put forth a **computational criterion for truth**: a statement is true if there exists a terminating, verifiable computation that accepts it (and no such computation refutes it). This criterion is executed through the BDI system for all the mathematics we build upon it. It ensures that the concept of truth remains tightly linked to evidence and repeatability. It demystifies truth by equating it with "what a computer could check given enough time/memory." This might sound like a reduction, but it actually empowers us – it means whenever you claim something is true, you are implicitly committing to there being a way to demonstrate it mechanically. Mathematics becomes a collaborative game between humans and machines to find these demonstrations.

1.6 Proof as Executable Trace and Formal Verification

Just as we redefined truth, we now redefine **proof**. In the traditional paradigm, a proof is a sequence of statements in a formal language, each following from previous ones by a rule of inference, starting from axioms and ending in the theorem. It is typically linear (or at best tree-structured, but usually presented linearly) and intended to be read by humans for validation. In our paradigm, a proof is an **executable artifact** – essentially a program or a structured graph that, when run, **produces the theorem from the premises**. The closest analogy is a computer program that transforms an initial state (encapsulating axioms and assumptions) into a final state that encodes the conclusion. Each step of the program corresponds to applying a rule or performing a valid operation in the appropriate DSL.

We already hinted that proof and program become one. Let's flesh this out. Suppose we want to prove a theorem in group theory: *If G is a finite group of odd order, then G is solvable* (the Feit-Thompson theorem). A traditional proof is hundreds of pages of group-theoretic arguments. A machine-oriented proof might involve a large sequence of deduction steps broken into lemmas – indeed this theorem was machine-verified in Coq, producing a proof term of enormous size ⁶. That proof term in Coq is essentially a lambda-term (a functional program) whose type is the theorem statement, by Curry-Howard. When Coq checks it, it is essentially executing that proof program to see that it yields the result. In BDI, we would similarly encode all those steps as a giant graph (with possibly subgraphs for lemmas). The **execution trace** of that graph is the proof – one can log each step (each BDI node firing with inputs and producing outputs) and that log is the evidence of correctness. Because the BDI nodes are designed to implement sound logical or computational steps, the whole execution's correctness is guaranteed if it terminates properly without contradiction. The *logged trace* can be checked independently by a simpler program that just knows how to verify each type of node execution (like a bytecode verifier). This is analogous to a proof checker going through each inference in a written proof.

Crucially, proof as execution means **proof becomes dynamic** rather than static. For example, if you have a probabilistic proof or an interactive proof, those fit well: they might involve random choices or interaction with an oracle, but as long as at the end you can verify the transcript by execution, it counts. Our system could even incorporate interactive proofs (where a prover and verifier exchange messages) by considering the entire interaction as one big distributed computation that results in the verifier being convinced. The log of that interaction would be the proof artifact. In fact, this resonates with modern complexity theory where they characterize various proof systems (NP, IP, PCP, etc.) – we could potentially realize those in BDI to let certain proofs be very efficient to check (e.g., a probabilistically checkable proof could be a thing where

the BDI verification randomly samples certain parts of the trace, etc., although implementing that securely is advanced).

The **verification** of a proof in our framework is multi-layered, as mentioned: does it run to completion? Does each step adhere to rules? Optionally, does it correspond to an external derivation or a trusted certificate?. We can formalize these. Each BDINode that is part of a proof has an `operation` field that might be something like `INFER_X` for some rule X (for instance, an operation type for applying a group theory axiom, or an `ARITH_ADD` that adds two numbers as part of checking some algebraic identity). The BDI execution engine can be instrumented to check *locally* that each `INFER_X` node's outputs truly follow from its inputs according to rule X. If all nodes pass this local check and the graph connectivity is acyclic and well-founded, the overall graph is a valid proof. This is akin to a kernel in a proof assistant checking each low-level step – except here the “steps” can be as fine as machine instructions if needed.

The optional ProofTags come into play if a node's correctness depends on an external theory. For example, if our BDI has a node type that says “By Fermat's Last Theorem, no solutions for this exponent”, we don't want to bake all of Wiles' proof into the BDI. Instead, we treat it as an oracle step with a ProofTag that refers to an external proof (like a DOI of a formal proof in Isabelle/HOL, say). The `META_VERIFY_PROOF` operation could then be used by the BDI engine when in audit mode: it might not expand Wiles' proof fully (because that's huge), but it can verify a cryptographic signature or hash that attests “someone checked this”. This way, our system can leverage prior verified results without re-proving them from scratch every time, while still not simply trusting them blindly – the trust is transferred via a secure hash or certificate. This is similar to how a modern proof assistant might accept a previously proven theorem as a lemma (under the assumption the theorem was proved in that system or another trusted system).

Because proofs are programs, we can apply **optimization** to proofs. A long proof might have unnecessary steps, which is like dead code. The BDI optimizer could remove or rearrange steps (as long as the dependency graph remains valid) to shorten the trace – effectively finding a shorter proof. This is analogous to proof compression techniques in logic (like cut-elimination is a kind of optimization, removing detours). Also, proofs could be **reused**: a proof graph, being a graph, could share subgraphs among multiple results. If two theorems have similar lemmas, the lemmas can be proved once and reused by pointer/reference. Traditional written proofs also do this by citing lemmas, but here it's literal reuse of the execution graph. If lemma L is proven, any other proof graph can just call L's subgraph as a subroutine (like a function call via `CALL` operation). This encourages a highly modular, component-based buildup of mathematical knowledge – much like software libraries. In fact, one can envision a *Mathematics Standard Library* of proven facts in BDI form, which can be linked into new proofs so that one doesn't reprove basic facts. This is analogous to the libraries in theorem provers (like the Compendium of Continuous Mathematics for Coq, etc.), but in our case the library is directly in executable form, not just a human-readable theorem database.

Another benefit is **animation of proofs**: since a proof can actually run, we can witness the process. For instance, a combinatorial proof that some configuration exists might actually construct an example during execution. This blurs the line between “proof by existence” and “algorithm to find it”. In classical math, non-constructive proofs give existence without a method. In our system, any existence proof is by necessity constructive (or else how did the program verify the existence? It must have found or built something). So proof execution often will output not just “true” but possibly a witness. We could design proof operations that yield witnesses as part of their postconditions. This turns every proven theorem of the form “there exists X such that P(X)” into an effective algorithm to find such an X (or at least as effective as the proof makes it).

The concept of **ledgery** that we have mentioned a few times refers to recording proofs in an immutable log (like a blockchain or hash chain). This would ensure that the history of proofs is preserved and can be audited. It's an extrinsic addition: even if one doesn't trust our system fully, one could independently hash each step of a proof trace and compare it to the ledger to make sure it hasn't been tampered with. This might be beyond pure math and into the realm of securing knowledge (imagine an open repository of all verified results, where any user can verify a hash to confirm they have the same proof content). Such a ledger could also serve as a priority record (who proved what when) or as a dataset for machine learning (mining patterns from known proofs to conjecture new lemmas). The synergy of formal proof and modern tech like distributed ledgers is speculative but intriguing.

In sum, by treating proofs as executable, **verifiable traces**, we align the practice of mathematics with the practice of debugging and verifying programs. Proofs become shareable, checkable digital objects. Mathematics becomes a **verification science** – each theorem is like a program specification, and a proof is the verification of that specification. This can dramatically increase confidence in results (no more subtle errors hidden in papers, as sometimes happens, because a machine check would catch them), and it opens up new ways for collaboration (since proofs are data, they can be transmitted, replayed, adapted).

We should note that this doesn't necessarily make proving easy – it's not a magic bullet to find proofs. But it ensures that once a proof is found (by human insight, AI, or brute force), it's automatically in a form that others (or other programs) can use reliably. It also means that an AI theorem prover can work hand in hand with a checker: the AI generates candidate proof graphs, and the checker verifies them, rejecting invalid ones. This tight feedback can even guide the AI (similar to how AlphaGo was trained by self-play with feedback, an AI prover could be trained by attempting proofs and seeing if the checker accepts them, gradually learning strategies that succeed).

By firmly embedding formal verification at the heart of mathematics, we essentially future-proof mathematical knowledge for an era of powerful AI and computer assistance. It means we can trust machines to explore mathematics because we have a rigorous way to validate their discoveries. This synergy between *formal verification* and *mathematical creativity* is a key motivation for the framework and will be explored in later chapters when we discuss intelligent systems.

1.7 The Binary Decomposition Interface (BDI): A Semantic Execution Fabric

Throughout this chapter, we have frequently referenced the **Binary Decomposition Interface (BDI)** as the enabling technology for our vision. We will now outline the role and design of BDI in bridging high-level mathematical DSLs with low-level binary execution. In short, BDI is proposed as a *unified, semantic execution fabric* for all computations – from mathematical proofs to AI algorithms – grounded in binary operations. It is the layer where Machine Epistemology concretely meets computer architecture.

Traditional computing systems have a steep abstraction hierarchy: high-level languages → intermediate representations (IRs) → assembly → binary machine code, often losing semantic information at each step. In contrast, BDI is designed to preserve *semantic richness* all the way down to execution. Think of BDI as a kind of universal assembly language that is *typed and self-documenting*. A BDI "program" is not a linear sequence of instructions but a **graph** $G = (V, E)$ where each node $v \in V$ is a structured operation (BDINode) and edges E represent data and control flow between operations. This graph is directly

executable by a BDI Virtual Machine (BDIVM), which can interpret the graph or compile it just-in-time to actual hardware instructions.

What makes BDI special is that each node carries **metadata linking it to high-level semantics**. For example, if a node represents adding two numbers, it will have metadata saying "this comes from a $++$ operation in Peano arithmetic DSL" or "this corresponds to an axiom of commutativity used here." The node structure, as we saw, includes fields like `operation` (the kind of operation), typed input/output ports, and `metadata_handle` which can reference things like source code location in the original DSL, proof annotations, etc.. It might also include a `region_id` indicating on what part of memory or which hardware region it should execute – this is important for performance (e.g. distinguishing CPU vs GPU memory).

By having this rich structure, BDI nodes are effectively **self-aware instructions** – they know what they are intended to represent, not just how to execute. This is crucial for verification: a node might be an `ASSERT` operation, which doesn't do computation but asserts a condition must be true; the BDIVM can check it at runtime. Or a node might be `RESOLVE_DSL` which takes a high-level construct (like a symbolic function) and dynamically links it to a BDI subgraph implementing it. This level of semantic fidelity is far beyond standard CPU instructions, which have no clue if an ADD is adding integers, floats, or parts of a data structure – they just add bits. BDI's nodes, by contrast, are tagged with type information and intent, which allows the system to reason about them at runtime or compile-time in ways compilers and debuggers currently struggle to.

The **graph model** also means that control flow is explicit as edges (no program counter except as an emergent property of following control edges), and data flow is explicit (edges carry results, no hidden registers unless represented as such). This is similar to dataflow architectures or synchronous circuits, but BDI can represent both dataflow and control flow networks. For instance, a loop is represented by a control edge going back to an earlier node (making a cycle, which is executed until a condition node breaks it). Conditionals (if-then-else) are nodes with a boolean input that gate which branch subgraph executes next, akin to a phi node in SSA form or a mux in hardware.

Now, how does BDI facilitate what we want for Machine Epistemology? It does so by being the **common exchange layer** between different DSLs and the hardware. A high-level proof written in, say, a proof assistant DSL would be compiled into a BDI graph. A machine learning algorithm written in PyTorch (a DSL for tensor operations) could also compile into a BDI graph. If both are in BDI, they can interoperate or at least coexist – for example, the proof could call a neural network as an oracle and verify its output, or a neural net training loop could incorporate a formally proven update step. BDI provides a common representation so that symbolic and numeric, deterministic and probabilistic computations all share a runtime. This addresses the *fragmentation* issue where different systems (CPU vs GPU, or theorem prover vs Python) have different languages and we cannot easily compose them. With BDI, the composition is by graph linking.

Another major feature is **verifiability**. Because BDI retains semantics, one could perform formal verification directly on BDI graphs. One could prove that a BDI graph implementation of some algorithm meets a high-level spec by induction on the graph structure, or by invariant checking. In fact, the BDI nodes include operations explicitly for verification, like `META_VERIFY_PROOF` which we discussed, or `META_ASSERT` which can enforce an invariant in the graph. These aren't normal operations in execution (they might only run in debug mode) but they serve to embed logical checks within the graph. This is like having **built-in**

Hoare logic: instead of treating asserts as comments, the BDIVM can actually monitor them. If an assert fails, it indicates a bug in the algorithm or an unsatisfied precondition, etc. Therefore, any BDI program can carry its own specifications and even proofs (as subgraphs) and can be set to self-verify at runtime. This is an implementation of the concept of **proof-carrying code** within the execution fabric itself.

From a hardware perspective, BDI is binary – ultimately everything is encoded in bits – but it is not tied to a particular machine word size or architecture. It could, for example, dispatch parts of the graph to different processors (some nodes to CPU, some to GPU, some to special accelerators) because the `region_id` and `hardware_hints` metadata can guide that. One can imagine a future where a single binary (the BDI graph serialized) runs efficiently on a heterogeneous computing cluster, optimizing itself. Because the semantics are intact, even such distribution would not alter the correctness: it's like the difference between a mathematical proof and the strategy to verify it – the proof is the same, whether one checks it by hand or in parallel or with different tools.

Graph vs. IR vs. ISA: It's worth comparing BDI to known representations. LLVM IR is a popular intermediate representation for compilers, but it's low-level (close to assembly) and not designed to carry full source-level semantics. It loses, for instance, the difference between a + used for integer addition versus one used for vector addition in source – it would have different instructions but doesn't know their conceptual link. BDI would keep them unified with type tags. Also, LLVM IR isn't meant to be directly executed normally, whereas BDI is designed for direct execution (like a VM). On the other hand, something like Python's bytecode is executable but too high-level and dynamic for verification. BDI aims to hit the sweet spot: low-level enough to be efficient and close to hardware, high-level enough to be understood and reasoned about.

Comparison to circuits: BDI graph can be seen as similar to an electronic circuit diagram. In a circuit, logic gates (nodes) are connected by wires (edges). Circuits can compute anything that a program can (given enough gates or if you allow cycles for state). Indeed, Shannon's insight was mapping Boolean algebra to circuits. BDI is like a *reconfigurable circuit* that can include high-level components (like an adder as one node, rather than thousands of gate nodes). You could implement a BDI VM in hardware directly, executing node by node. In fact, one might compile BDI to an FPGA, effectively turning the program into a hardware instance for speed. Because region and hardware hints exist, BDI is prepared to exploit parallelism and special hardware.

In summary, the Binary Decomposition Interface is the **keystone** that enables our philosophical framework to be realized in practice. It ensures that our lofty talk of “everything is binary and verifiable” doesn't remain an ideal but becomes a concrete software (and perhaps hardware) reality. BDI provides the structure to represent any mathematical or computational process in a uniform way. It elevates the binary level from an “implementation detail” to a **first-class semantic layer**. In doing so, it promises to drastically reduce the gap between thought and execution: one unified representation that both humans (at least via tools) and machines can work with.

The rest of this book (and research) will delve deeper into BDI's design (Chapter 4 as referenced in snippet appears to focus on BDI) and how it can be used to build complex intelligent systems (Chapters 2 and 3 on computation and cognition, respectively, set the stage for why BDI is needed). For now, we have an initial understanding: BDI is our *machine epistemology workbench* where all our theories are meant to be realized.

1.8 Operational Semantics and Metadata of BDI

To appreciate the robustness of BDI, let us consider its **operational semantics** (how programs execute on it) and the role of various metadata structures (which carry proofs, types, and other annotations). The operational semantics of BDI can be thought of in two layers: (1) **Execution Model**: how the control and data flow in the graph is scheduled and carried out, and (2) **State Transformations**: what each node does to the machine state (registers, memory, etc.) when it executes. We ensure that both layers are defined rigorously so that one can reason formally about a BDI program's behavior.

Execution Model: A BDI graph, as noted, may have complex structure including potential cycles (for loops) and multiple entry/exit points (for functions or subgraphs). We impose that each graph has designated start and end nodes (`META_START` and `META_END` in the `OperationType` enum), which mark where execution begins and where it should conclude. When a BDI program is run, the BDIVM will place a token at `META_START` and then follow control flow edges deterministically (or nondeterministically if specified, say for parallel threads) to traverse the graph. Control flow edges define a partial order of execution. If multiple control outputs emanate (like a branch node), the one whose condition is true will be followed, or both if it's parallel fork. Data flow edges ensure that when an operation node executes, it has the needed inputs from predecessors – the VM might stall an operation until all its data inputs are ready (like dataflow execution). This is similar to how out-of-order CPUs or dataflow architectures work, avoiding executing something before its operands are computed.

The BDI execution thus naturally supports **parallelism**: independent parts of the graph can execute concurrently. This matches how proofs can be checked in parallel or how independent sub-computations can run simultaneously. The semantics guarantees the same result as a sequential topological order execution (since the graph defines dependencies clearly). We can formally prove a *confluence* property: any execution order respecting the dependencies yields the same final state (this is crucial for determinism). If a graph is acyclic, it's purely functional; if it has cycles, it implies state, which BDI represents through special memory and state nodes (like a `MEM_STORE` followed by a later `MEM_LOAD` which gets that stored value creates a cycle conceptually but with a memory versioning semantics). The Region and Port constructs allow precise tracking of what state is being read/written.

State is also structured: Memory in BDI is segmented into regions, which can represent different types of memory (heap, stack, GPU memory, disk, etc.) or different abstract data spaces (like a region representing the variables of a mathematical model). By tagging memory operations with region identifiers, the semantics can ensure no illegal interactions (like you can't mix integers from region A with floats from region B if not allowed, etc.). This provides a built-in memory safety model. It also aids formal verification because aliasing is controlled (two pointers with different region IDs are definitely not aliasing each other).

Operational semantics of nodes: Each `BDIOperationType` has a rule for how it transforms inputs to outputs. For example, an `ARITH_ADD` node will consume two integer inputs and produce their sum as an output (modulo some bit-width if specified). A `BRANCH` node will consume a boolean and based on it, direct control to one of two control outputs (if present). A `CALL` node will invoke a subgraph: its semantics might be expanding inlined or jumping to that subgraph's start and later returning. We keep these semantics clear so that one could, if needed, reduce a BDI graph to a logical formula or to a behavior trace. The presence of `META_*` operations is purely for meta semantics: `META_COMMENT` does nothing (just

holds text), `META_NOP` does nothing (just a placeholder), etc.. Those exist to carry documentation or to pad graphs.

One particularly important set of operations are those for **Proof and Verification** (like `META_ASSERT`, `META_VERIFY_PROOF`, `PROOF_TAG` etc.). Their semantics are not about computing a number but about checking conditions in the *logical* realm. For instance, `META_ASSERT` with an input bool will terminate the program or flag an error if that bool is false (thus enforcing an invariant) – otherwise it just passes control along. `META_VERIFY_PROOF` might take as input a ProofTag (some representation of a proof certificate) and output true if the certificate matches and false or error if not. The exact details depend on the design, but the key is that the semantics of these nodes tie into *external logic checking*. We can, for formalism, treat them as axiomatic: assume `META_VERIFY_PROOF` is sound (if it says the proof is valid, then in an ideal sense the theorem holds because an external checker signed off). Alternatively, one could not use `META_VERIFY_PROOF` at all and instead fully expand proofs in BDI for purely internal verification – which is more self-contained but potentially inefficient.

Metadata structures: BDI offloads a lot of information into metadata to keep the core graph lean. Metadata might include: the original source code (in some DSL) that led to this node (helpful for debugging or round-tripping proofs to human-readable form), names of variables or theorems, types of variables beyond simple bit types (like a node could have a metadata saying "this 32-bit value represents a float according to IEEE754"), hardware mapping suggestions (e.g., this operation is memory-bound so send it to DMA engine), and cryptographic hashes of expected behavior (ProofTags as discussed). This metadata is not necessarily used during normal execution, but it can be used by tools: e.g., an optimizer might use type info to simplify something, or a debugger might use source info to show a trace to the user in familiar terms. The separation of metadata ensures we don't clutter the execution with checks for things that are not needed every time. But the **presence** of metadata is what distinguishes BDI from a typical IR – it remains available so that at any point we can *interpret the graph at multiple levels of abstraction*.

For example, if a BDI program fails an assert, the system can use metadata to print "Assertion failed: $x > 0$ was false at line 50 of user's code", because it knows what that assert meant in source terms. Or if we want to prove a property of the BDI graph itself, metadata might contain an invariant that was suggested by a human, which a model-checking tool can use as a hint.

The **graph model with metadata** also naturally supports *proof-carrying code* more directly than classical PCC. In classical PCC, the proof is an external entity attached to code. In BDI, the proof can live in metadata or as separate subgraphs that are linked via ProofTags. For instance, one could have a BDI subgraph that is the formal proof in a DSL like Lean, and a ProofTag that is a hash of that subgraph. Then in the main computation, a `META_VERIFY_PROOF` node uses that tag to confirm correctness before allowing some critical operation (say a type-cast or an array bounds omission that is justified by the proof) to proceed. This way, the code literally *contains its proof* in a checkable way. An analogy: in Rust programming language, the type system proves at compile time certain safety, but those proofs are erased at runtime. In BDI, we might keep such proofs around as metadata and optionally check them at runtime if desired (for extra safety or for audit).

Domain-Specific optimization and DSL integration: There is an operation type mentioned as `RESOLVE_DSL`. This likely refers to an operation that translates or dispatches a part of the graph that's written in a high-level DSL into lower-level BDI subgraph. For example, if we had a symbolic polynomial object and we want to perform an integration, `RESOLVE_DSL` might trigger the system to replace that

with actual arithmetic nodes that compute the result. It's somewhat like a just-in-time compilation of domain-specific constructs. The semantics here might be more complex (it might involve looking up a library or rule set at runtime). But crucially, because BDI allows *self-modifying graphs* or *meta-operations*, we can implement things like reflective reasoning (a proof that generates subproofs algorithmically).

Integration with external systems: Sometimes BDI will need to call out to external code (perhaps legacy code or OS services). This can be handled by having special node types like `EXTERN_CALL` with a specification of what it does, or by treating external systems as separate regions with known semantics (like reading a sensor or user input could be a non-deterministic input node). The design should ensure these are sandboxed so as not to violate verifiability. Possibly such calls require a proof obligation (the user must assert that the external call meets certain spec, akin to assume-guarantee reasoning).

At this point, we see that the **BDI acts as a broad umbrella** under which many concepts from programming languages, logic, and hardware meet: it is simultaneously an IR (like SSA), a VM (like JVM), and an ontology for knowledge (with metadata linking to formal semantics). It embodies the idea "Logic is instruction, Memory is topology, Proof is execution trace, Learning is graph transformation" – each of those slogans indicates how BDI treats traditionally separate aspects as unified. Logic is literally implemented by instructions that carry logical semantics. Memory structure (how data is laid out) becomes part of the execution topology as graph regions. Proofs we already covered as traces. And learning (for AI) could be seen as modifying the graph (like adding edges or adjusting weights which are stored in memory nodes) – something BDI can do and verify if those transformations meet some criteria.

The operational semantics of BDI is defined in such a way that **if a BDI graph encodes a mathematical statement and its proof, executing the graph effectively is checking the proof**. This is the realization of "make computation and inference the same" – so we no longer distinguish a program from a proof script, it's all just BDI graphs. This rigorous alignment of semantics ensures that the high-level correctness proven in a DSL is maintained down to the bits flipping in the hardware. In IEEE style terms, BDI provides *end-to-end semantic preservation*.

To conclude this technical deep dive: the BDI's execution model and metadata architecture are carefully designed to satisfy the demands of **Machine Epistemology**: - They guarantee that every operation is grounded in binary (ontological primacy of 0/1 states). - They preserve the meaning of higher-level constructs (semantic fidelity of knowledge). - They allow direct execution (no large gaps between reasoning and running). - They facilitate verification and linking of proofs (intrinsic verifiability). - They are universal across domains, encouraging unification of disparate systems of knowledge in one substrate.

This is an ambitious design, but as we will see in later chapters, it offers immense potential payoff: more **verifiable, efficient, and intelligent** computational systems built on a substrate that "understands the meaning behind the bits". BDI is, in essence, the machine infrastructure that embodies our machine epistemology principles.

1.9 Verifiability, Proof-Carrying Code, and Symbolic Compression in Action

To cement the concepts discussed, it is useful to envision how all these pieces come together in practice on a concrete example. Consider a scenario: we want to create a high-assurance software component that

implements a mathematical function – for instance, a cryptographic algorithm that relies on number theory (say, elliptic curve operations), and we want to ensure its correctness and security properties formally. Using our framework, we would approach this as follows:

1. **DSL formulation:** We first describe the cryptographic algorithm in a high-level mathematical DSL – for example, a DSL that has constructs for finite field arithmetic, elliptic curve points, etc. In this DSL, we might state a specification: e.g., "Compute kP (scalar multiplication on the curve) correctly for all inputs k, P ," and also prove properties like the output is still on the curve, the algorithm runs in sublinear time in k 's bit-length, it doesn't leak k through timing (a security property), etc. This DSL-level proof might use algebraic reasoning, perhaps relying on group theory lemmas.
2. **Compilation to BDI:** We compile the high-level description into a BDI graph. This means every operation, like adding two field elements, becomes a sequence of low-level BDI nodes (bit-level additions with mod reduction, etc.). The proof that we had at the high level is also compiled or linked. If the proof was done in a proof assistant, we attach it via ProofTags to relevant parts of the code. Alternatively, we may have proven some correctness properties by code annotations (like loop invariants) which become `META_ASSERT` nodes in BDI at those points.
3. **Proof-Carrying Execution:** When we deploy this BDI program, it essentially carries a proof of its own correctness. For example, there might be a top-level `META_VERIFY_PROOF` that checks a certificate that "this implementation conforms to the math spec and has no timing side-channels" – something which could be proved by static analysis and attached. The end user or the host system, upon loading this BDI program, runs the verifier part, which goes through the proof embedded and confirms everything. This is an instance of **proof-carrying code**: the code can be allowed to run only after the proof is validated. Unlike original PCC which was often outside the code, here it's intertwined, but conceptually similar.
4. **Verifiable execution:** As the code runs (say it's doing kP for some input k), the BDI VM will also keep an eye on the assertions. If a bug were present such that an assertion can fail (meaning some proved invariant wasn't actually true due to a mistake), the system would catch it at runtime (assuming we didn't catch it in proof, which ideally we would – but consider maybe a hardware fault flips a bit, then an assert might catch that the result is off, acting as a safety net). Also, because of region management, we have built-in safety against buffer overflows, etc. So the execution is robust.
5. **Symbolic compression and introspection:** Suppose we want to optimize the algorithm or understand it better. The BDI graph is semantically rich, so a tool can read it and notice patterns. For instance, it might identify that the sequence of operations doing field multiplication corresponds to a known polynomial multiplication algorithm and replace it with a more efficient one (like using a different method or hardware accelerator). This optimization is done preserving semantics thanks to the type and intent metadata (similar to how a compiler uses types, but here even more powerful with proofs guiding it maybe). This shows *symbolic compression*: the high-level concept of field multiplication is still recognizable in the low-level BDI code via metadata, allowing reasoning at that higher abstraction even during optimization.
6. **Audit and trace:** If later an auditor or another program wants to ensure everything happened correctly (say during a security audit), they can replay the BDI execution trace in a sandbox, check the ledger of hashes to ensure it wasn't tampered with, and confirm the proof tags correspond to

known theorems or certified results. This gives an unprecedented level of assurance: they see exactly which theorem (with hash X) was used at which step, and they can cross-verify theorem X in some database of formal proofs.

This example encapsulates what we gain: **trustworthiness** (via proofs and verification), **interoperability** of high-level math and low-level code (via DSL to BDI pipeline), **efficiency** (the code can be optimized and run on hardware directly), and **transparency** (one can inspect and audit every piece). It also shows how our initial philosophical tenets percolate down: the truth of the crypto algorithm's correctness is tied to an executable proof. The proof is just as much a part of the deliverable as the program itself, blurring the line between theory and implementation.

Looking at a more scientific example: Suppose a researcher conjectures a certain combinatorial property and they write a program to search for counterexamples. In our framework, they could incorporate the conjecture as a `META_ASSERT` inside their search algorithm. The search (a BDI graph) will then either find a counterexample or finish without violating the assert. If it finishes without violation, that execution trace itself is a proof of the conjecture up to the searched bound. If they generalize it inductively, they may turn it into a full proof with the help of a theorem prover. The same environment (BDI) can handle the brute-force search and the inductive proof, unifying experimental and theoretical mathematics.

This synergy is powerful for **symbolic AI** as well: an AI agent using this system could reason about its own strategies (proofs about why its decision will be safe), and execute them with guarantees. For instance, a planning AI could prove that "under these conditions, this plan achieves the goal" and carry that proof with the plan execution code so that a monitor can verify the plan is on track. This is essentially what you would want in a safety-critical AI (like a self-driving car proving it will not violate certain safety conditions given assumptions).

Thus, the interplay of **verifiability** and **symbolic compression** in BDI is not an academic exercise; it is meant to solve real problems in building reliable, understandable complex systems. Each higher concept (like a loop invariant, or a mathematical law) is compressed into machine-enforceable snippets, which guide the execution and provide hooks for checking. It's like having a very smart runtime that knows about the logic of the code, not just the mechanics.

To wrap up the chapter, let's reflect on how this approach constitutes a *foundational redefinition* of mathematics and computation: - Mathematics is no longer detached from implementation; it is explicitly a layered construction on binary computation. - Computation is no longer a rote symbol shuffling; it is elevated to carry meaning and even *know* what it's doing relative to human concepts. - The binary substrate is not a limitation but the enabler of absolute rigor: because bits are unambiguous and machines don't get tired, we can demand 100% proof and verification, as opposed to the traditional mathematical practice which relies on social verification (peer review, reputation, etc.). Our approach aims for a world where every claim can be checked by anyone (or any machine) to arbitrary detail.

By grounding everything in binary state transformations, we are essentially embracing the motto "**It from bit**" in a practical way. We assert that every "it" (mathematical truth, program, intelligent behavior) comes from "bits" – and we have built the scaffolding to trace that emergence step by step, with no mysterious gaps. This gives us confidence that as systems scale, we can still trust them, because they can be *made to explain themselves* all the way down to flipping bits, and those explanations can be verified.

1.10 Conclusion: Executable Foundations and the Road Ahead

In this chapter, we embarked on a journey to reformulate the foundations of mathematics and intelligent systems through the lens of Machine Epistemology. We established the philosophical stance that **verifiable computability** is the core of meaning and truth in mathematics. Distilling that philosophy, we re-examined fundamental mathematical constructs – numbers, sets, logic, functions – and redefined them in terms of executable operations on a binary substrate. We argued that mathematics can be viewed as a network of **Domain-Specific Languages**, each a deliberate compression of underlying binary processes, and that by preserving the connection to those processes, we can achieve unprecedented rigor and automation in mathematical reasoning.

A key theme has been the convergence of proof and program. We no longer see a proof as a static argument on paper, but as a dynamic, checkable computation – a stance encapsulated by the Curry-Howard correspondence and realized in practice by our approach ³. In doing so, we align with ongoing trends in formal methods and theorem proving, but push them further by integrating them with the fabric of computation itself (via BDI).

The **Binary Decomposition Interface** was introduced as the pivotal technology enabling this integration. BDI serves as the "enzyme" that catalyzes the reaction between high-level abstract mathematics and low-level machine code, ensuring no information is lost in translation. It is, effectively, an attempt at the *Holy Grail* of computing: a system that is both highly abstract and fully concrete, combining the strengths of human reasoning (abstraction, generalization) with the strengths of machine execution (speed, precision, verification). BDI elevates the binary bit from an implementation detail to an **ontological primitive** of our knowledge representation, fulfilling the vision that the most secure and sound knowledge is that which can be run and tested at the simplest level.

By rooting our epistemology in the binary substrate, we gain a kind of **universality** reminiscent of the Church-Turing thesis: any process of reasoning or computation, no matter how complex, in principle can be embedded in this framework. The challenge, of course, is one of scalability and usability – writing raw BDI graphs is like writing assembly, not feasible for humans at scale. But that's where the layered DSL approach shines: one works in whatever high-level language is suitable (be it set theory, Python, or a new AI logic), and the BDI ensures it all maps down to a common, verifiable core. It's akin to how in software engineering we can code in Python but have confidence that down the line, C and assembly and microcode will faithfully carry out our intent, except we add the twist that at each level we keep *proofs of correctness* in the loop.

The *philosophical implications* are significant. We are, in a sense, proposing a new answer to the ancient question "What is mathematics?" – mathematics is **not** about an ethereal realm of ideal objects knowable by pure thought, nor just a formal game with symbols, nor merely a useful language. It is the emergent structure of what can be **reliably done** by an information processing system. It is the ultimate compression of executable knowledge. This perspective harmonizes with the physicalist view (Wheeler's "it from bit", the idea that physical reality itself might be information at bottom) and with the pragmatist view (that knowledge is what we can do or verify). It does alter how we view mathematical truth: no longer an absolute in the Platonic sense, but a conditional – true *when* the requisite computation exists. However, since that condition can often be met for the statements of interest (especially in finitely checkable domains), this does not limit mathematics, it grounds it.

From the standpoint of **computability and complexity**, our framework provides a very clear measuring stick. A theorem is not just true/false, but also associated with the *complexity of its proof (execution)*. This opens the door to classifying mathematical problems by the resources needed to verify them – an area that intersects with complexity theory. We could ask: is this conjecture provable with a short (polynomial-size) proof or only a gigantic one? These become well-defined questions in the BDI context and could guide research (maybe some statements are inherently complex to verify, akin to lower bounds in complexity theory).

As we proceed in the book, Chapter 2 will build on this foundation to discuss how general computation (beyond pure math proofs) can be seen as structured binary processes, introducing concepts like entropy, complexity, and adaptation in computational systems. Chapter 3 will apply these ideas to *intelligent systems*, proposing an operational definition of intelligence in terms of recursive, verifiable computation over binary memory. By Chapter 4, we return in depth to the BDI, detailing its architecture, comparing it to existing systems (like how it differs from LLVM, JVM, etc. as glimpsed in the snippet), and showcasing examples. Subsequent chapters likely tackle applications, optimizations, and perhaps a new language (Chimera was mentioned as a paradigm on BDI ⁷).

To conclude this chapter, we reaffirm the grand vision: a future where every piece of knowledge – whether a mathematical theorem, a piece of software, or an AI’s strategy – comes with a **certificate of its own validity** that any computer can check. A future where mathematics and science are fully digital, not just in presentation but in essence, enabling automation without sacrificing rigor. In this future, errors and bugs have nowhere to hide, as they would manifest as failed verifications at some level. Creativity will still be needed (to find proofs or invent new DSLs for new domains), but once something is created, its verification becomes routine and trustable.

This is a radical departure from the status quo, but examples in limited domains (formal verification of hardware, use of proof assistants in verifying critical software, etc.) already show it’s possible and beneficial ⁴. Our contribution is to unify these efforts under one conceptual roof and push them further into the mainstream of how we define mathematics and build intelligent systems. If successful, the payoff is a **tower of knowledge with solid foundations** – from the bit up to human-level ideas – a tower we can build as high as we want, confident that it will not collapse because every joint and beam has been checked.

The journey has just begun, but with the foundational principles and tools laid out in this chapter, we have a map for the road ahead. Each subsequent chapter will add more structure and detail to this framework. By the end, we hope to convince the reader that Machine Epistemology is not only a philosophically appealing stance but a practical and necessary one for the next era of mathematics and intelligent systems – an era where human creativity and machine reliability join forces to expand the realm of knowledge.

References:

1. J. A. Wheeler, *Information, Physics, Quantum: The Search for Links*, Proc. III Int. Symp. Foundations of Quantum Mechanics, 1989 – *Introduced the “it from bit” concept emphasizing binary distinctions as fundamental to physical reality.*
2. C. E. Shannon, *A Mathematical Theory of Communication*, Bell System Tech. Journal, 1948 – *Established the bit as the fundamental unit of information and linked Boolean logic to electronic circuits* ¹.

3. G. C. Necula, *Proof-Carrying Code*, POPL 1997 – *Describes attaching formal proofs to code so that machine can verify safety properties before execution.*
4. M. Mernik, J. Heering, A. M. Sloane, *When and How to Develop Domain-Specific Languages*, ACM Comp. Surveys 37(4), 2005 – *Overview of DSL design, defining DSLs as specialized programming languages for particular domains* ⁵ .
5. D. A. Pierce, *Types and Programming Languages*, MIT Press, 2002 – *Explains the Curry-Howard correspondence where proofs correspond to programs and propositions to types* ³ .
6. G. Gonthier et al., *A Machine-Checked Proof of the Odd Order Theorem*, Interactive Theorem Proving 2013 – *Formalized the Feit-Thompson theorem (odd order of finite groups implies solvable) in Coq, demonstrating that complex math can be verified by machine* ⁴ .
7. **BDI Design Document**, 2025 – *(Provided in the materials)* Detailed description of the Binary Decomposition Interface, including the BDINode structure, operation set, and philosophy behind a semantic binary-executable graph.
8. **Binary Mathematics Manuscript (Ch. 1 Draft)**, T. Mohammed, 2025 – *(Provided)* Preliminary draft articulating the concept of Machine Epistemology and foundational redefinitions (some text integrated in this rewrite).
9. Anonymous, *QED Manifesto*, 1994 – *Called for formalizing all of mathematics in a computer-verifiable way* ² , foreshadowing the goals we pursue with modern tools.
10. W. A. Howard, *The Formulae-as-Types Notion of Construction*, 1980 – *Manuscript articulating the Curry-Howard isomorphism, identifying proofs with programs, a key theoretical pillar of our approach* ³ .

¹ How Claude Shannon Invented the Future | Quanta Magazine

<https://www.quantamagazine.org/how-claude-shannons-information-theory-invented-the-future-20201222/>

² QED manifesto - Wikipedia

https://en.wikipedia.org/wiki/QED_manifesto

³ Proofs are Programs - YouTube

<https://www.youtube.com/watch?v=AGnTnbR1sSg>

⁴ Six-year journey leads to proof of Feit-Thompson Theorem - Phys.org

<https://phys.org/news/2012-10-six-year-journey-proof-feit-thompson-theorem.html>

⁵ When and how to develop domain-specific languages

<https://dl.acm.org/doi/10.1145/1118890.1118892>

⁶ A Machine-Checked Proof of the Odd Order Theorem - SpringerLink

https://link.springer.com/chapter/10.1007/978-3-642-39634-2_14

⁷ 2025-04-27-Binary Mathematics and Intelligent Systems by Tariq Mohammed Chapters 1-10.pdf

<file:///file-UBwwFk7FpYVMod3u9msGdJ>