

Chapter 4: The Binary Decomposition Interface (BDI) – A Foundational Computational Substrate

Abstract – This chapter introduces the Binary Decomposition Interface (BDI), the core computational substrate underpinning the Binary Mathematics framework and realizing the principles of Machine Epistemology. The BDI is presented not merely as an intermediate representation (IR) or compiler stage, but as a universal, graph-based computational fabric that directly bridges high-level symbolic reasoning (expressed in diverse domain-specific languages) with verifiable, machine-level execution. We detail its formal structure based on typed binary nodes and semantic graphs, its functional roles transcending traditional compilation pipelines, and its unique capabilities in enabling proof-carrying code, hardware-aware semantic optimization, and introspectable execution. BDI provides the necessary architecture for grounding abstract mathematical and intelligent processes in traceable, executable binary dynamics, unifying execution semantics, proof evidence, and optimization in a single substrate.

4.1 Rationale: The Need for a Semantic Execution Fabric

Modern computing stacks suffer from a **semantic gap** between abstract logic and machine execution. As discussed in Chapter 1's Machine Epistemology, all verifiable knowledge (mathematical theorems, program logic, learned models) must ultimately be traceable to operations on a binary substrate. However, as we move from high-level specifications to physical hardware, semantic richness is progressively lost. Traditional compilation pipelines (Source → AST → IR → Assembly → Machine Code) tend to strip away high-level context and intent at each lowering stage. For example, an LLVM IR or bytecode often retains only rudimentary type info and operations, losing the original program's invariants or proofs of correctness. This gap poses serious challenges:

- **Fragmentation:** Different hardware (CPUs, GPUs, FPGAs, accelerators) require distinct IRs or instruction sets. Code is re-targeted and rewritten for each platform, hindering portability and unified reasoning about program behavior.
- **Opacity:** Low-level binaries and assembly lack traceability to source logic. Verifying that machine code upholds the intent of the source (or mathematical spec) is non-trivial – it often relies on external formal methods or extensive testing. Even formally verified compilers like CompCert aim to prove the translation correctness, but the artifact (machine code) itself remains an opaque sequence of bits requiring trust in the toolchain ¹ ².
- **Lack of Integration:** Intelligent systems rely on concepts like adaptive feedback loops, entropy tracking, and proof traces (Chapter 3), but current hardware and software interfaces provide no first-class support for these. Traditional ISA and IR designs don't natively capture “entropy of a variable” or “proof of safety” as part of the program representation.
- **Safety and Verifiability:** Ensuring safety properties (memory safety, type safety) at low level is difficult. Approaches like proof-carrying code (PCC) [1] and typed assembly [2] were proposed to attach formal proofs or type constraints to low-level code, but these are add-ons to a fundamentally unsafe substrate. There is no ubiquitous execution format today that *intrinsically* carries its verification evidence.

In summary, there is a critical gap between abstract DSL specifications rich in semantics and their efficient, verifiable execution on hardware. The BDI is designed to overcome these limitations by providing a unified, **semantic-preserving**, verifiable, and executable substrate that sits directly above raw binary hardware but below high-level languages. It serves as a universal translation and execution layer where meaning is preserved rather than lost.

4.2 Philosophical Foundations and Core Principles

The design of BDI flows from the philosophical stance of Machine Epistemology and the needs identified above. BDI treats the **binary distinction (0/1)** as the ontological primitive of computation – not just an implementation detail, but the very canvas on which all symbolic structures must ultimately be painted. This leads to several core tenets guiding BDI's architecture:

1. **Binary Primacy:** All data, code, and metadata in BDI reduce to binary representations. At the lowest level, every node and connection is encoded in bits. This echoes the idea of a *typed assembly language* [2] pushed to its extreme: even complex mathematical objects or proofs are ultimately realized in binary form within BDI. By enforcing binary grounding, we ensure that nothing is beyond machine verification – if it exists in BDI, it exists as bits that can be checked and executed.
2. **Semantic Fidelity:** Unlike conventional IRs which sacrifice high-level information, BDI aims to preserve semantic intent and structure from the source throughout the pipeline. The *meaning* of a construct (e.g., a loop invariant, a matrix dimension, or a logical precondition) is not thrown away during compilation but embedded as metadata on BDI nodes. This principle aligns with the motivation behind multi-level IR frameworks like MLIR, which recognize the value of carrying richer information in IR. BDI pushes further: semantic tags and provenance are first-class citizens, enabling the system to reason about “why” a node exists, not just “how” it operates.
3. **Executability Mandate:** Every BDI structure is directly executable or interpretable. A BDI graph isn't an abstract math object – it *is* a program that can run on a virtual machine or be compiled to silicon. This contrasts with proof objects in systems like Coq, which are mathematical certificates not intended to execute. In BDI, even proofs and assertions are part of an executable graph. Thus, BDI serves as both the *specification* and the *implementation*. As a consequence, there is no final translation step that loses confidence – the BDI graph you verify is the thing you run.
4. **Verifiability and Proof-Carrying Code:** BDI is built to be *self-verifying*. Inspired by the concept of Proof-Carrying Code [1], each part of a BDI graph can carry proof metadata that certifies certain properties (safety, functional correctness with respect to a spec, etc.). Instead of trusting external proofs after the fact, the proof evidence travels *with* the code. For example, a node performing a memory store may carry a proof obligation (and ideally a discharged proof) that an index is in bounds. The BDI runtime or compiler can automatically check these proofs, and even refuse to execute nodes that don't meet their proof requirements. This approach unifies formal verification with execution – the line between code and proof blurs.
5. **Compositional Graph Structure:** BDI represents programs as graphs, which are inherently compositional. Smaller graphs (or subgraphs) can be composed to form larger ones without losing their correctness properties. This reflects how complex systems are built from simpler components. It also allows **multiple DSLs** or paradigms to interoperate on one substrate – e.g., a logic proof graph can connect to a numeric computation graph – because all are just BDI subgraphs. The interface between them is binary and graph-based, eliminating impedance mismatch between different runtime systems.

6. **Hardware Awareness and Tunability:** While BDI is architecture-agnostic in semantics, it is designed to be aware of hardware characteristics. Each node can carry *hints* or requirements for hardware execution (e.g., “this is a GPU-friendly operation” or “prefer to execute near memory bank X”). This allows BDI to act as a layer where high-level algorithm knowledge meets low-level performance tuning. Unlike Java bytecode or WebAssembly which deliberately abstract away the hardware, BDI can expose hooks for optimization. For instance, a BDI node could specify that it should reside in an FPGA’s DSP block versus general logic. The BDI is not locked to one machine model; instead, it can adapt to *any* machine by appropriate mapping of its regions and operations.

In essence, BDI is conceived as the layer where **abstract logic meets executable dynamics, and where proofs meet programs**. It treats binary not just as bits of data but as the fundamental medium of meaning. By upholding semantic richness down to the binary level, BDI aims to fulfill the Machine Epistemology mandate: to make knowledge truly executable and verifiable all the way down.

4.3 Formal Structure: BDI Graphs, Nodes, and Typed Metadata

At its core, a BDI program is not text but a **typed, binary-executable graph** $G = (V, E)$. The vertices V are *BDI nodes* representing operations or data, and the edges E represent typed dependencies (flow of control, data, or memory effects) between nodes. Figure 4.1 illustrates a simple BDI graph structure with a few nodes and edges.




Figure 4.1: A simple BDI graph fragment. Each circle or box is a **BDINode** (with an identifier, operation type, and a designated region for execution or storage). Blue arrows indicate data dependencies: for example, two **LOAD** nodes (loading values from Region R0) feed an **ADD** node, whose result then flows into a **STORE** node writing to Region R1. Each node carries a *RegionID* indicating the logical memory or compute region it targets (here R0 might be an input memory region, R1 an output region, and the **ADD** executes in a CPU region). Control flow edges (not shown for simplicity) can enforce sequencing when data dependencies alone are not sufficient (e.g., to order side effects). This graph representation makes both data flow and control flow explicit as edges in a unified structure.

BDI Nodes (V): Every node $v \in V$ is a structured binary object encapsulating a wealth of information beyond a typical instruction. Conceptually (simplified for exposition), we can think of a BDI Node as having the following fields:

- **NodeID:** A unique identifier (e.g., a UUID or content-derived hash). This ID allows any node to be referenced unambiguously, enabling cross-graph references, caching of compiled code, and proof attachments. The NodeID also serves auditing and traceability – one can log execution of NodeID 42 and tie it back to a specific high-level construct.
- **OperationType:** An opcode or tag indicating what the node does. BDI defines a rich set of operations, ranging from typical arithmetic (`BIN_ADD` , `BIN_MUL` etc.) and memory ops (`MEM_LOAD` , `MEM_STORE`), to control flow (`CTRL_BRANCH_IF`), to meta-operations (`DSL_RESOLVE` for embedding DSL logic, `META_ASSERT` for assertions, `VERIFY_PROOF` for checking proofs, etc.). The operation type places the node in the broad taxonomy of behaviors known to the BDI execution engine.
- **Typed Payload:** Optional immediate data or parameters for the operation, with an associated type tag (the *BDIType*). This can include numeric literals, vector lengths, addresses, or even inline data like small lookup tables. By tagging the payload with a type, BDI ensures that even immediate values are interpreted in a known, typed way (for instance, a payload might be an 8-bit unsigned int vs a 32-bit float, explicitly).
- **Data Inputs / Outputs:** Rather than each node containing opaque operand slots, BDI nodes list their input dependencies and output ports explicitly. An input is a reference to another node's output (by NodeID and an output port index). Outputs are described by a type and semantic name. These replace the role of registers or temporary variables in conventional IRs. The collection of data edges forms a dataflow graph: if node B's input comes from node A's output, an edge $A \rightarrow B$ (data) exists. These edges are strongly typed – the output type of A must match the expected input type of B, enforced by BDI's type system at graph construction time.
- **Control Inputs / Outputs:** These list the control-flow predecessors and successors of the node, defining possible execution order (like basic block connections in a CFG). For example, a `BRANCH_IF` node might have two control outputs (true and false branches). Control edges ensure that even in the absence of data dependencies, the execution order is well-defined (important for operations with side-effects or for sequencing things like I/O).
- **Metadata Handle:** Crucially, each node carries a reference (handle or pointer) to a rich metadata object capturing high-level semantic information. This **SemanticMetadata** can include the originating DSL code or AST node, variable names, documentation strings, source line numbers, logical assertions or specifications associated with this operation, and proof objects or cryptographic hashes of proofs. The metadata is effectively an extensible slot to hang arbitrary *non-executable* information that should travel with the node. Unlike debug info in typical compilers (which can be stripped without affecting execution), BDI's metadata can influence optimization and validation (e.g., a proof in metadata might be required for the node to execute).
- **Region ID:** A label indicating the logical memory or compute region that this node is associated with. "Region" in BDI generalizes the concept of memory spaces and execution contexts. Examples of regions include a specific memory segment (heap vs stack, or a particular allocator pool), a CPU cache level, a GPU's shared memory or SM, an FPGA block, or even an abstract region like "persistent storage." By tagging nodes with regions, BDI can enforce and optimize locality – e.g., a `MEM_LOAD` tagged with RegionID 5 will be executed with respect to that region's memory space. Regions also help reason about non-interference (operations in distinct memory regions might not conflict) and capabilities (a node might only access a region if it has permission).

- **Hardware Hints (optional):** Hints for lower-level scheduling such as preferred execution unit, expected latency, alignment constraints, vectorization info, etc. Rather than baking these into the operation semantics, they live in metadata or a structured hint field. For instance, a matrix multiplication node might carry a hint “this is GPU-amenable” or “tile this operation for caches.” These do not change the functional semantics but guide the compiler or VM for performance.
- **ISA Binding (optional):** For very low-level nodes or when interfacing with existing machine code, a node can carry an explicit binding to a sequence of machine instructions (for a specific architecture) that implements this operation. This is similar in spirit to having an assembly snippet for a high-level operation. If present, it enables direct verification against hardware: one could formally prove that the bound instruction sequence correctly implements the node’s abstract operation (as done in verified compilers [6]), or use it to quickly realize this node on that hardware.
- **Execution Properties:** Flags or enum describing the node’s execution behavior – e.g., is it deterministic and pure (no side effects)? Is it idempotent or reversible? Does it potentially throw exceptions or traps? Does it read/write memory (and if so, which region)? These properties assist in analyses and optimizations. For example, knowing a node is pure and deterministic means it can be safely recomputed or memoized; knowing it has no side-effects means it can be reordered more freely.

In summary, a single BDINode packs what in a conventional system would be spread across multiple layers: an opcode (from an instruction set), operand connectivity (from a dataflow graph), debugging annotations, type information, and even formal proof links. This **rich node structure** is what allows BDI to be semantic-preserving and self-verifying. The design is akin to a *typed assembly instruction with attached proof certificate and hardware map* – a concept foreshadowed by work on PCC and TAL [1][2], but integrated here into a unified executable model.

BDI Edges (E): Edges represent the connections between nodes. As noted, we distinguish different types of edges implicitly by what they connect: data edges (from an output port of one node to an input port of another), control edges (from a node’s control output to another’s control input), and memory dependency edges. The latter deserve a note – because memory operations in different nodes might refer to the same region, there can be dependencies to enforce correct ordering (for example, two writes to the same memory region should be ordered even if their data flows don’t intersect). BDI can represent memory ordering constraints either as special control edges or via a memory model in the metadata that analyzes overlapping regions. In essence, edges ensure that the *graph as a whole captures all necessary relationships* (data requirements, ordering, resource usage) for correct execution. The resulting structure isn’t just a control-flow graph or a dataflow graph, but a **multi-layered semantic graph** that captures *what* is being computed, *how* it connects to other computations, *where* it should execute (regions), and *why* it is correct (metadata/proofs) ³ ⁴. It is this holistic nature that differentiates BDI from simpler IR graphs.

Illustrative Example: Suppose we have a high-level expression $y = a + b * c$. In a traditional SSA-based IR like LLVM, this might become two instructions: `%t1 = mul i32 %b, %c` and `%y = add i32 %a, %t1`. The high-level context (that this was computing a linear expression) is lost; we just have two operations in sequence. In BDI, we would create a small graph: nodes for the constants/inputs `a`, `b`, `c`; a `MUL` node with inputs from `b` and `c`; and an `ADD` node with inputs from `a` and the output of the `MUL`; finally an output node `y` capturing the result. This forms a DAG representing the expression. Figure 4.2 shows a side-by-side comparison of such an IR vs BDI representation.

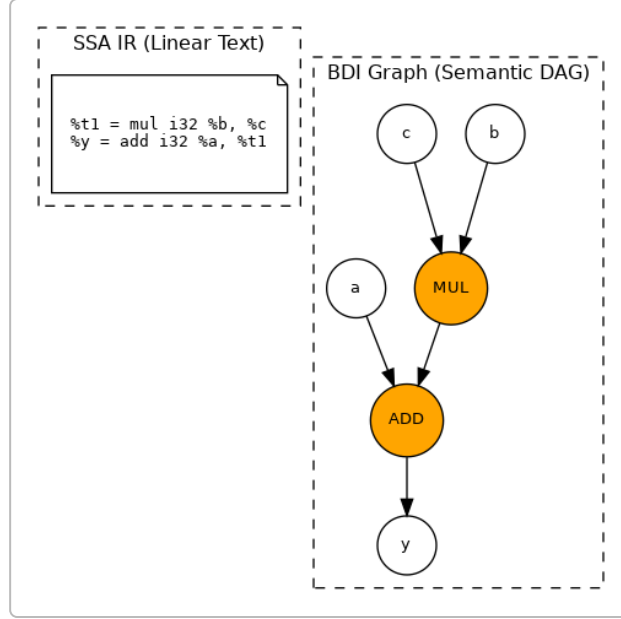


Figure 4.2: Traditional SSA IR vs BDI Graph. **Left: SSA-style linear IR** for `y = a + b * c` with two instructions (mul then add). **Right: BDI semantic DAG** for the same computation, where `a`, `b`, `c` are input nodes (white), the `MUL` and `ADD` are operation nodes (orange), and `y` is an output node. The BDI graph explicitly shows the data dependencies (arrows) and doesn't require a temporary variable name for the intermediate result. Moreover, additional metadata (not shown here) could be attached to the `MUL` or `ADD` node (for example, an indication that this operation came from a polynomial computation in the DSL, or a proof that `ADD` will not overflow if that was a DSL invariant). In the SSA text, such contextual information is absent. BDI's graph carries a richer story about the computation than the raw SSA does.

By structuring all computations as graphs of typed, metadata-rich nodes, BDI sets the stage for a substrate where high-level semantics and low-level executability co-exist. In the following sections, we explore how this structure is utilized as a multi-faceted engine – serving simultaneously as IR, as a virtual machine code, and as a formal proof container.

4.4 Functional Roles of the BDI Substrate

Because BDI integrates information typically scattered across different layers (source semantics, IR, optimization hints, and even proof artifacts), it can fulfill multiple roles that traditionally belong to separate tools. In a conventional stack, we have distinct components: compilers, linkers, VMs, optimizers, verifiers, etc. The BDI graph, by virtue of its completeness, blurs these boundaries. We highlight several key roles that BDI takes on:

1. Universal Semantic Translator: BDI acts as a common target into which diverse high-level formalisms can be translated *without losing meaning*. A compiler or translator for a DSL (be it a math language, a data science notebook, or a formal proof script) would emit a BDI graph that embodies the original semantics. For example, a theorem from Coq or Lean can be compiled into a BDI subgraph where each inference step is a node (with attached proof metadata), effectively creating an *executable proof*. A numerical algorithm written in Python could be translated into an equivalent BDI graph with each high-level operation (matrix multiply, say) broken down into precise binary operations annotated with the original intent (matrix

dimensions, etc.). The crucial point is that this translation is *semantic decomposition*: it preserves correctness-by-construction. The high-level concepts are resolved into BDI operations (via `DSL_RESOLVE` nodes and such) with provenance links. BDI becomes a universal *exchange format* for knowledge and computation, akin to an assembly language for everything from algorithms to proofs.

2. Compiler Internal Representation and Backend: BDI can serve as both **the IR and the backend target** in a compilation pipeline. A traditional compiler might lower a language to an IR (like LLVM IR) then further to machine code. With BDI, the DSL can lower directly to a BDI graph. Once in BDI form, there are multiple execution paths: - **Interpreted Execution:** A BDI Virtual Machine can interpret the graph directly, executing node by node. This is useful for rapid prototyping or when just-in-time feedback is needed (similar to how Java bytecode is often interpreted before hotspot compilation). - **Ahead-of-Time (AOT) Compilation:** The BDI graph can be *directly* compiled to machine code for a target architecture, using the node's `ISA_Binding` or via pattern-matching subgraphs to machine idioms. Since BDI nodes can carry machine-specific bindings or be lowered in groups, one could bypass any textual assembly stage and generate binaries straight from the graph. This is an “IR-free” or rather “single-IR” compilation – no separate lowering to another IR like RTL or JVM bytecode is needed. In essence, BDI is the final IR. - **Optional Lowering to Existing IRs/HDLs:** In scenarios where we want to leverage existing compilers or hardware synthesis tools, a BDI graph (or parts of it) can be lowered to standard representations like LLVM IR, SPIR-V, or even Verilog/VHDL ⁵. For instance, a kernel subgraph heavily oriented to GPU could be translated to SPIR-V [4] to use existing GPU drivers, if needed. Or a fixed-point arithmetic subgraph might be converted to Verilog to implement on FPGA. The ability to export BDI to other formats provides compatibility and a migration path, but ideally BDI-to-hardware will be handled through its own toolchain in the long run.

In this role, BDI also replaces the linker: BDI graphs can be merged simply by taking the union of nodes and hooking up edges accordingly (with appropriate renaming of NodeIDs to avoid collisions). The usual linking issues (symbol resolution, relocation) are simplified since NodeIDs can be globally unique hashes, and external references can be explicit edges to library BDI graphs.

3. Runtime Execution Environment: A BDI Virtual Machine (BDIVM) serves as the execution engine for BDI graphs. Conceptually, BDIVM plays a role analogous to a CPU or a language runtime (like the JVM), but at a higher abstraction level. When running a BDI graph, the BDIVM is responsible for scheduling node execution (respecting control/data dependencies), managing memory regions, and interfacing with actual hardware resources. Execution in this environment enables features like: - **Live Introspection:** Because everything is in a graph, the VM can pause and inspect the state of any node, the values on edges, or the metadata. This is like a supercharged debugger where one can ask not only “what is the value in this register?” but “what is the logical condition this node was meant to check?” (accessible via metadata). One could step through a proof carried by the code while it executes, with the system checking each proof step. - **Dynamic Modification (Adaptive Learning):** The BDIVM can permit controlled modification of the graph at runtime. This is key for intelligent systems that learn or adapt. For example, nodes can be added or rewired on the fly as new learning rules activate (much like modifying a neural network's structure during training). Because changes are themselves BDI operations (e.g., a special `GRAPH_REWRITE` meta-operation), they can be verified or rolled back if unsafe. This enables *self-modifying code* in a disciplined, verifiable manner – a crucial capability for systems that evolve (Chapter 3's adaptive graphs). - **Consistent Cross-Platform Model:** The BDIVM provides the same logical execution model regardless of underlying hardware. Whether the physical machine is x86-64 or a many-core RISC-V or a distributed cloud, the BDI execution semantics don't change – only the mapping (via regions and hardware hints) does. This is analogous to how the JVM

gives Java portability, but here the execution model is lower-level and more deterministic, with less abstraction of hardware.

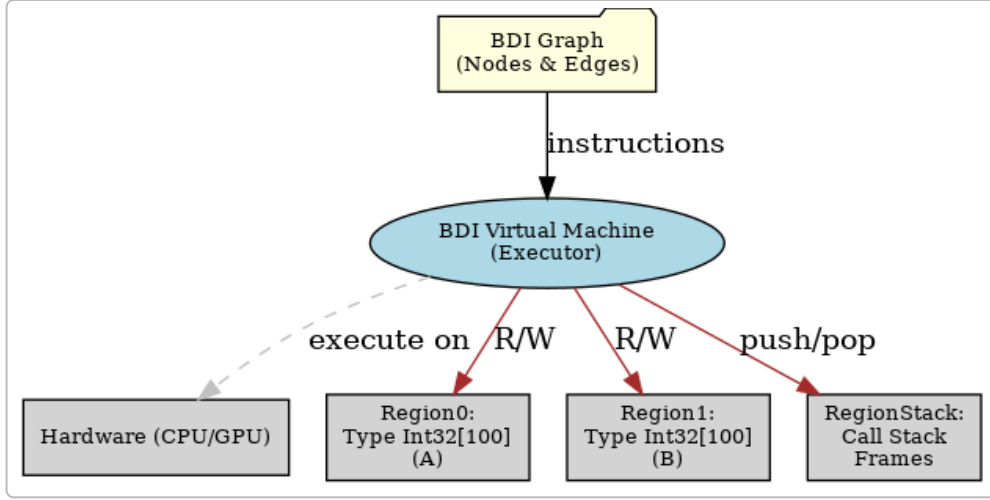



Figure 4.3: BDI execution model integrating memory regions. **The BDIVM (center)** reads “instructions” in the form of the BDI graph nodes and orchestrates execution. It interacts with various **Memory Regions (bottom)**, here Region0 and Region1 (which could represent separate memory spaces or different data sets) and a special RegionStack for call stack frames. Each memory region is a typed, bounded space – for instance, Region0 might be an array of 100 int32 values. The BDIVM ensures that memory operations (`MEM_LOAD/STORE`) to these regions are checked against bounds and types. The BDIVM can also leverage hardware: in this schematic, a dashed arrow indicates that it may dispatch operations to actual CPU/GPU hardware as needed, guided by hardware hints. The BDI Graph (top) provides the BDIVM with the next node(s) to execute (akin to “fetching instructions”, though the flow is graph-driven rather than a flat sequence). This diagram emphasizes how execution, memory, and hardware interface all converge in BDI: the graph tells *what* to do, the regions define *where* data lives, and the BDIVM bridges it to *how* to run on real silicon.

4. ISA Semantic Modeler: BDI can also function as a framework for modeling and reasoning about hardware instruction sets themselves. By representing each machine instruction as a BDI graph pattern or node template, BDI can serve as a *formal semantics layer for ISAs*. For example, an add instruction in x86 could be modeled as a BDI node `X86_ADD` with certain ExecutionProperties (it’s arithmetic, affects flags, etc.) and an ISA_Binding that ties it to opcode 0x01 of x86. Why is this useful? It allows us to perform *typed assembly programming* [2] natively in BDI – one could write low-level code with the assurance of BDI’s type and proof checks. It also permits formal verification at the ISA level: sequences of actual machine instructions can be lifted into BDI graphs and checked against higher-level specs. This realizes the vision of PCC [1] and CompCert’s verified compilation [6] in a more flexible way: instead of proving a monolithic compiler correct, we prove individual instruction rules correct by modeling them in BDI, then ensure our BDI graph transformations preserve correctness. In other words, BDI can be the meeting ground where high-level verified algorithms and low-level machine instructions are compared and verified.

5. Proof-Carrying Code Framework: Perhaps most ambitiously, BDI serves as a built-in proof-carrying code framework. The term PCC typically refers to an architecture where code producers supply a machine-checkable proof that the code respects certain properties (such as memory safety) [1]. In BDI, proof carrying is native. Each node or group of nodes can carry proof artifacts (via `ProofTag`) or metadata linking to

proofs in systems like Coq/Lean). The *entire graph* can thus be seen as a proof-carrying entity, where the safety policy is embedded in the node definitions themselves (e.g., type rules, region access rules) and the “certificate” is the collection of proof metadata showing those rules are followed. The BDIVM can be envisioned as not just an executor but also a proof checker: as it executes, it can verify any proof obligations on the fly. For instance, if a node is marked with an assertion “x must be positive here” and a proof is provided, the VM can check that proof (or simply validate the condition at runtime if it’s a dynamic check) and halt if it fails. BDI graphs are thereby *self-verifying*: they contain the evidence of their own correctness. This dramatically reduces the trusted computing base for safety: we no longer need to *trust* the compiler to emit safe code; the code *shows* its safety, and the runtime can enforce it. Recent research has moved toward attaching proofs to compiler IR (e.g., to LLVM bitcode [7]), and BDI builds on these ideas, generalizing them to an entire system representation.



*Figure 4.4: A proof-carrying BDI node (conceptual). The orange BDI node represents a memory store operation `STORE arr[i]`. A green **Proof node** (or metadata) above it encodes the assertion and proof that `i < len(arr)` (the index is within bounds). The dashed green arrow indicates that the proof *verifies* the safe execution of the store. In practice, the proof node could be a `META_ASSERT` operation with a proof attached, or simply an annotation on the store node. During execution, the BDIVM would verify this condition (either by checking the proof or by runtime check if no static proof is provided). If the proof is valid, the store proceeds; if not, an exception or halt occurs. By structuring code this way, BDI ensures memory safety intrinsically. Contrast this with a traditional system where bounds checks are either inserted as separate runtime tests or proved offline – here the proof lives *inside* the program graph. This approach can be extended to other properties (overflow avoidance, functional correctness of a routine with respect to a spec, etc.), making BDI a vehicle for *certified programs*.*

Through these roles, BDI essentially unifies what today are disparate concerns: high-level language semantics, low-level code generation, runtime execution, and formal verification. The next sections will delve deeper into specific technical aspects – particularly how execution and memory are managed in BDI, and a comparative analysis of BDI against state-of-the-art technologies that address parts of this puzzle.

4.5 Execution Semantics and Memory Architecture in BDI

Execution Model: A BDI graph is conceptually executed by “flowing” through the directed graph respecting dependencies. Unlike a linear instruction stream, multiple nodes could, in principle, execute in parallel if they have no dependencies between them (much like instructions in dataflow architectures). The BDI execution semantics can be viewed as a *scheduling of the graph* such that for every node, all its control and data inputs have been produced before it executes. This is similar to the semantics of synchronous dataflow or classical static single assignment (SSA) form – except BDI also considers control edges and memory ordering as prerequisites.

- **Determinism and Concurrency:** By default, a BDI graph describes a deterministic computation (given the same inputs and initial memory state, it will produce the same outputs), because the dependencies enforce a partial order. However, concurrency can be naturally expressed: independent subgraphs can execute concurrently, and join via control nodes. BDI’s explicit control edges can model parallel forks, joins, and even nondeterministic merges if needed (though one would typically encapsulate nondeterminism via special nodes or external inputs).
- **Conditional and Iterative Execution:** Control flow structures are represented in BDI by control nodes (like `CTRL_BRANCH_IF`, `LOOP_START`, `LOOP_END`, etc.) and possibly by *graph cycles* for loops (a cycle in the graph indicates iterative execution). A loop in BDI might be represented by a back-edge for control and appropriate phi-like nodes for data that persists across iterations. The presence of cycles means the graph is not strictly acyclic in those cases; execution semantics then involve fixed-point or repeated firing until a condition changes. This brings BDI close to how hardware description languages (HDLs) represent circuits (which can have cycles but with understood feedback semantics).
- **Function Calls and Stack:** A function call can be represented by a special node (e.g., `CALL`) that has control outputs to the callee graph and an eventual control return back. BDI can manage call stacks using a dedicated memory region (RegionStack as shown earlier) or via special nodes that push/pop frames. Because BDI can represent recursive structures (graphs can reference subgraphs by NodeID or via metadata), recursion and reentrant calls are supported. The key is that each call invocation can be assigned a new RegionStack segment for its local variables, ensuring isolation and type safety of the stack frame.

Memory Regions and Data Safety: Memory in BDI is not a single, flat array of bytes as in conventional architectures. Instead, memory is partitioned into **regions**, each of which has an associated type and size (and potentially other attributes like alignment or location). For example, one region might represent an array of 100 integers, another a hash table of a certain structure, another a block of GPU memory with float16 tensors, etc. This resembles higher-level languages where objects have types and sizes, except here it’s at the substrate level.

- **Typed Pointers:** In BDI, when a node refers to memory (say a `MEM_LOAD` or `STORE`), it doesn’t use a raw pointer but rather a reference like (RegionID, index) or (RegionID, address) depending on abstraction level. Because the region knows its element type, BDI can infer the type of data being loaded. This effectively provides **spatial safety**: accessing beyond the bounds of a region can be detected (the index or address can be checked against region size), and **type safety**: one region can’t be erroneously interpreted as another type. This is in spirit similar to the goals of Typed Assembly Language [2], which extended type guarantees to the lowest level code. BDI’s approach makes such safety intrinsic. A buffer overflow in BDI would be an illegal memory edge going outside

a region's declared bounds – something that can be caught during graph construction or at runtime checks, preventing the unchecked behavior seen in standard machine code.

- **Region Capabilities:** We can associate permissions with regions (read-only, read-write, execute, etc.). A node that tries to write to a read-only region would violate a semantic rule, again catchable by verification. This aligns with modern safety features (like Rust's borrow checker or capability machines), but implemented as part of BDI's semantics.
- **Memory Dependencies:** If two nodes access the same memory region, BDI by default conservatively assumes a dependency (to maintain sequential consistency within that region). However, through metadata, one can specify finer grain information (like two loads from region R that target different indices have no write/write conflict). BDI could allow parallel reads, etc., but if there's a write involved, edges (or an intervening fence node) will enforce order. The memory model of BDI can thus be stricter (easier to reason about) than typical weak memory models of CPUs. One can, of course, relax it for performance by adding explicit unordered annotations if needed.
- **Interfacing with External Memory:** If BDI is running on real hardware, at some point it must map its logical regions to physical memory (RAM, GPU memory, etc.). The BDI runtime or compiler's *Hardware Abstraction Layer* (HAL, see Chapter 7) is responsible for allocating actual memory for each region and ensuring the BDI memory operations translate to actual loads/stores at the correct addresses ⁶ ⁷. This layer can also handle DMA transfers for regions that reside on accelerators, etc. The key point is that from the perspective of the BDI program, memory is segmented and typed – providing a structured view that aids verification and optimization.

Example – Safe Array Handling: In a traditional C program, an array access `arr[i]` compiles to something like: compute address = base + i*sizeof(elem), then load or store at that address. If `i` is out of bounds, this is a runtime error but nothing in the machine code prevents it. In BDI, one would model `arr` as a Region with length `N` and type `T`. The access is a node with inputs `(RegionID_arr, index=i)`. BDI semantics would dictate a check: `0 ≤ i < N`. This check could be an explicit `ASSERT` node connected to the `STORE`, or simply part of the `STORE` node's defined behavior (requiring a proof or runtime check). Thus, either at compile-time a proof is attached that `i` is in range, or at runtime the VM will perform the check. This way, a buffer overflow is caught or proven impossible. The cost is either a proof obligation or a runtime conditional, but given BDI's design, such checks can often be discharged by proofs (especially if the code came from a high-level spec that guaranteed the range).

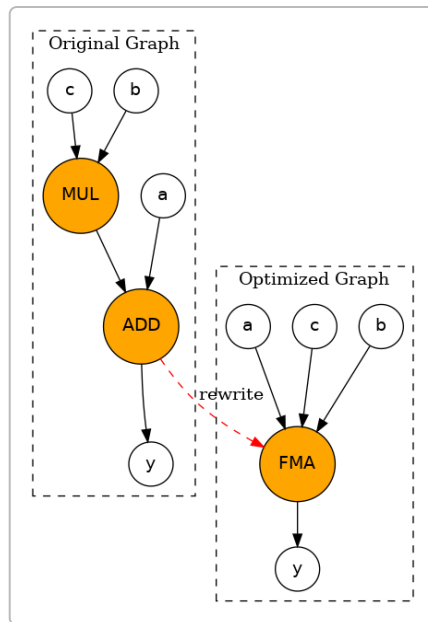
In practice, BDI's memory model resembles a safe systems programming model (like Java's or Rust's safety guarantees) but at the lower level: memory safety is a built-in property rather than an add-on. Research on PCC and TAL in the late 90s [1][2] showed it's feasible to enforce safety through types and proofs at low level; BDI incorporates those insights and extends them to general semantic properties.

Scheduling and Optimization: The BDI execution semantics leave room for various scheduling strategies. Because the graph explicitly delineates dependencies, a scheduler (statically in a compiler or dynamically in the VM) can choose an execution order or parallel assignment that respects the partial order. This is similar to how an out-of-order CPU or a parallelizing compiler works, except the "instruction" set here is much more expressive. BDI's rich metadata can assist scheduling: - For instance, nodes tagged to the same Region might ideally be executed close in time (to improve locality). - Nodes with hardware hints indicating "GPU"

might be batched and offloaded together. - A node marked as high-latency but with no data dependencies downstream could be executed earlier (pre-fetching, in effect).

The presence of control flow means BDI can't freely reorder across certain boundaries (just as one wouldn't move code across a branch in normal programs without care), but within a basic block-like region, dataflow gives flexibility. In essence, BDI enables a form of global dataflow optimization naturally: e.g., common subexpression elimination is trivial because one would naturally reuse the same node for a computation rather than have two identical sub-nodes. Similarly, dead code elimination is just removing nodes with no outputs used. Many classic optimizations become simple graph transformations on BDI.

Feedback-Driven Graph Rewriting: One powerful aspect of BDI's execution model is that it's amenable to *on-line optimization and learning*. The running system can monitor performance or accuracy and modify the graph accordingly. For example, if a certain subgraph is a bottleneck, the system might replace it with an optimized equivalent subgraph (perhaps using a different algorithm but proven to be equivalent). Or if this is an AI system, it might restructure part of its decision graph as it learns new rules. Because the graph is the program, such changes are effectively self-optimization. They can be done while preserving correctness if accompanied by proofs or by testing in a sandboxed way.



*Figure 4.5: Example of feedback-driven graph optimization. **Original graph (left)** computes $y = a + b * c$ using separate `MUL` and `ADD` nodes as in our earlier example. **Optimized graph (right)** replaces those with a single `FMA` (fused multiply-add) node that does the multiply and add in one step (if the hardware supports it). The red dashed arrow indicates a transformation: the subgraph with `MUL` and `ADD` is rewritten into the `FMA` node. Such a rewrite could be triggered by a peephole optimizer noticing the pattern, or by a machine learning system that predicts better performance with a fused operation. In BDI, this transformation would be performed by generating a new node (FMA) and redirecting the edges accordingly, then removing the obsolete nodes. If the transformation's correctness is not obvious, a proof can be attached (perhaps referencing that FMA on IEEE-754 floats yields the same result as separate operations, under certain conditions). This showcases how BDI unifies representation and optimization: the*

running program *is* a data structure that can be analytically optimized or learned from, rather than an opaque binary.

State and Entropy Considerations: BDI can model stateful computations (like a counter incrementing) by having a node that feeds its output back into itself through a region or loop construct (creating a feedback edge). Unlike pure dataflow which has trouble with state without specialized constructs, BDI's inclusion of regions (especially if a region stands for a register or memory cell) and control flow allows stateful loops, reactive systems, etc. We can imagine a BDI-based operating system (BDIOS) in which even the OS state (scheduler queues, etc.) is a BDI graph/region manipulated by BDI operations, making the whole system introspectable and verifiable.

In conclusion, the execution semantics of BDI marry the advantages of dataflow (explicit parallelism and clarity of dependencies) with the practical necessities of control flow and state. The memory model enforces safety and locality through structured regions. The result is an execution model that can be *as fast and low-level as C*, yet *as safe as a managed language*, and *as transparent as a formal model*. Next, we compare BDI to various existing technologies to highlight these points in context.

4.6 Comparative Analysis: BDI in the Landscape of Computing Models

To better appreciate BDI's design, we contrast it with several state-of-the-art representations and paradigms, highlighting differences and similarities. BDI draws inspiration from many of them but also diverges in critical ways.

4.6.1 BDI vs. Traditional Compiler IRs (e.g., LLVM IR) – Semantic Graph vs. Linear SSA: Traditional mid-level IRs like LLVM IR are text-based or in-memory representations of code organized in basic blocks and SSA form. They are primarily optimized for *compiler transformations* and close reflection of machine operations. In contrast, BDI is a graph-based representation capturing higher-level intent. A few key distinctions:

- **Representation:** LLVM IR is essentially linear text with labels and simple types (e.g., integers, floats) operating on an infinite register set (SSA values). It needs complex analyses to recover high-level info. BDI is fundamentally a *typed graph*. Nodes in BDI carry rich info (operation semantics, types, context) intrinsically. The graph form makes dataflow explicit without phi-nodes (merging happens through explicit nodes or structured loops) and allows naturally representing parallel flows. The BDI graph is more akin to an *extended dataflow graph with control*, whereas LLVM IR still has a sequential flavor with jump instructions forming a control flow graph ⁸ ⁹.
- **Semantic Fidelity:** As previously noted, much of the source semantics (e.g., variable names, loop invariants, etc.) is lost in LLVM IR. It's *possible* to attach metadata in LLVM (debug info, source mapping), but optimizations often ignore or even drop it. BDI, by design, maintains semantic metadata as a first-class part of the node. For example, if a calculation comes from a matrix multiplication in the DSL, the BDI nodes might know they're part of a matrix op, enabling domain-specific optimization that a generic compiler might miss. Another example: a proof that a loop is correct can be attached to the loop node in BDI, whereas in LLVM IR such a proof would have no standard place to reside. Essentially, BDI is *lossless w.r.t semantics* up to the detail the translator provides ¹⁰ ¹¹.

- **Verifiability:** Very few IRs have built-in verification beyond type checking. Efforts like translating IR into Coq to verify transformations exist (e.g., Vellvm for LLVM, and PCC for LLVM bitcode [7]), but those are external. BDI's approach is more integrated: by including proofs as part of the IR, the "correctness" of code can be checked *on the fly*. This is a step beyond even CompCert's approach [6], which proves the compiler that produces the IR is correct – BDI would allow checking properties of the program itself at runtime or load-time. In that sense, BDI provides a more *live* verification, closer to Necula's original PCC vision [1] but generalized.
- **Executability:** LLVM IR is not typically directly executed; it's meant to be compiled to machine code (though an interpreter exists for debugging, it's slow and not the primary use). BDI is explicitly meant to be directly executable by a VM. This makes BDI more akin to a **virtual instruction set** (like JVM bytecode or WebAssembly) but with higher-level semantics intact. One could deploy BDI graphs as program artifacts in a system, whereas one wouldn't deploy raw LLVM IR in production typically ^{9 12}.
- **Hardware Mapping:** LLVM IR is nominally platform-independent but in practice, it assumes a generic von Neumann architecture with flat memory and will get specialized during lowering. BDI is more abstract in that it doesn't assume a specific hardware layout – memory is via regions, compute can be anywhere. Yet it's also more concrete in that it can specify hardware preferences. It's as if BDI operates one level above LLVM IR, but then can dip one level below it as well by modeling the ISA.

In summary, BDI's semantic graph approach stands out against traditional IR by preserving more information, enabling built-in proofs, and being directly runnable. It addresses what Chris Lattner identified as limitations of SSA-based IRs (hence the development of MLIR) – indeed "SSA isn't enough" [8] in terms of capturing multi-level program representation – BDI takes that philosophy further by unifying even runtime and verification aspects.

4.6.2 BDI vs. High-Level Virtual Machines & Bytecode (JVM, .NET CLR, Python VM): High-level VMs like the Java Virtual Machine and .NET CLR define portable bytecode instruction sets that abstract away the underlying machine. They have features like object-oriented operations, garbage collection, and runtime type checking. How does BDI compare?

- **Abstraction Level:** The JVM and CLR are higher-level than BDI in many respects – they have opcodes for creating objects, calling methods, performing type casts, etc. BDI operates at a lower level of abstraction (closer to hardware) but carries high-level *metadata*. For example, the JVM has no concept of the cache or specific memory regions – memory is one big heap managed by GC. BDI, on the other hand, might explicitly distinguish different memory regions (heap, stack, GPU memory). So BDI is simultaneously lower-level in operation (no built-in notion of classes or objects) yet higher-level in the sense of being semantically rich (it could carry an object's schema in metadata even if it doesn't have a "new object" opcode per se) ^{13 14}.
- **Type System:** The JVM's type system includes classes, interfaces, etc., and enforces type safety at bytecode verification time. BDI's type system is simpler at the core (it deals with binary data types, numeric types, etc.), but can reflect high-level types through metadata or composite nodes. One advantage of BDI is that it can make *value-dependent* types more explicit (e.g., type indexed by a size known at runtime, which high-level VMs usually can't express). Additionally, BDI's approach to types

is extensible (via DSLs), whereas JVM/CLR have a fixed set of types they understand at the VM level

15 .

- **Semantic/Proof integration:** JVM and CLR verify certain properties (e.g., no operand stack underflow, type correctness of calls) before running bytecode – but they do not carry explicit proofs of arbitrary properties. BDI, as discussed, can integrate proofs for things like algorithmic correctness or domain-specific invariants, which is beyond the scope of mainstream VMs. In essence, BDI could ensure not just type safety (like JVM does) but *full functional correctness* of critical routines if proofs are attached.
- **Execution Model:** Both JVM and CLR use a stack-based execution model (WASM too, discussed next). BDI is graph-based – more akin to how modern JIT compilers internally turn bytecode into SSA graphs before native code. One could say BDI externalizes the graph structure that a JIT would internally use. This means BDI can naturally expose parallelism (multiple nodes ready to execute) whereas something like the JVM requires explicit threading to achieve parallel execution.
- **Performance and Optimization:** The high-level VMs rely on JIT compilation to achieve performance, turning bytecode into optimized machine code at runtime. BDI could leverage AOT or JIT as well, but interestingly, it also allows *semantic optimization* that a JVM can't do easily. For example, the JVM doesn't know that a particular sequence of bytecodes is actually computing, say, a matrix multiplication, so it can't replace it with a call to an optimized BLAS library without sophisticated pattern recognition. In BDI, that knowledge could be preserved and an optimizer could directly swap the subgraph for an optimized implementation (perhaps provided by a domain-specific library in BDI form, or through hardware instructions).

In short, compared to high-level VMs, BDI is *leaner* in baked-in features but *richer* in carried information. It's like a minimalist core (binary ops, control, memory) with an extensible semantic layer, whereas JVM/CLR have a fixed feature-rich core but little awareness of semantics beyond what the language front-end encoded as bytecode. BDI also aligns with some of the goals of these VMs – portability, safety – but extends safety to new dimensions (proofs, formal verification).

4.6.3 BDI vs. WebAssembly (Wasm): WebAssembly [5] is a modern, safe, portable binary format, often considered a bytecode for the web. It's lower-level than JVM bytecode (closer to an assembly for a virtual RISC machine), yet designed for safety and performance. A comparison with Wasm is illuminating:

- **Binary Format:** Both BDI and Wasm are binary formats intended for efficient execution. However, Wasm is inherently a *linear* stack-machine code (with structured control flow). A Wasm module defines functions which consist of instructions in a sequence (with block structures for control). BDI is a graph format and does not rely on an implicit operand stack – dependencies are explicit edges. This makes BDI more amenable to direct analysis (no need to simulate a stack to see how data flows).
- **Safety and Verification:** Wasm was designed to be memory-safe (through linear memory with bounds checking) and type-safe (validated before execution). It achieves PCC-like safety without requiring proofs by severely limiting what code can do (no arbitrary pointer arithmetic without bounds checks inserted by the engine). BDI similarly enforces memory and type safety through its design (typed regions, etc.). The difference is BDI can carry proofs for higher-level properties,

whereas Wasm's verification is hardwired to basic safety only. That said, both share the philosophy of *explicitly verifying code before/during execution* – Wasm just doesn't allow code that fails its validation to run, analogous to BDI refusing to execute a graph that violates type/region rules or fails a proof check.

- **Hardware Abstraction:** WebAssembly is a virtual ISA for a generic 32-bit stack machine with some modern features (like local variables, linear memory, etc.). It abstracts away the actual hardware (endianness, word size differences). BDI similarly abstracts hardware, but it can intentionally include hardware-specific nodes or hints when needed. BDI could target Wasm as a backend (i.e., generate Wasm code from a BDI graph), but doing so might lose some of BDI's richer metadata unless it's somehow encoded in custom sections (which wouldn't affect execution).
- **Performance:** Wasm is designed to be a compact, easily JITable representation, and indeed current Wasm engines reach near-native speed by optimizing the code. BDI's performance would depend on its implementation – interpreted BDI might be slower than interpreted Wasm because BDI does more (e.g., proof checking). But a compiled BDI could match or exceed performance if it leverages the extra info (e.g., it might auto-vectorize a calculation because the graph form makes it clearer, or skip redundant checks that it knows are proven safe). The BDI vs Wasm trade-off is between *generality* (BDI can represent things beyond a typical ISA, like proof logic) and *simplicity* (Wasm is intentionally minimal to ease engine implementation).
- **Use Cases:** Wasm is mainly for deploying applications on the web in a safe manner. BDI's scope is broader: it is meant as a foundation for any computational system, including OS kernels, AI models, etc. In some sense, BDI could serve as an *architectural description language*, whereas Wasm is firmly an *executable format*. We can imagine writing an entire OS in BDI (with device drivers, etc., all in graph form) – doing that in WebAssembly would be possible but awkward (especially for low-level device access, which BDI handles via its HAL and region abstractions).

In summary, WebAssembly and BDI both champion portability and safety, but BDI injects a heavy dose of semantic awareness and flexibility, at the cost of being more complex and ambitious. WebAssembly is a subset of what BDI can represent (one could encode a Wasm module as a BDI graph with mostly straightforward opcodes, but the converse is not true without losing information).

4.6.4 BDI vs. Formal Proof Assistants (Coq/Lean) and Proof Objects: Systems like Coq and Lean produce *proof terms* or objects that certify the validity of propositions. These proofs are typically not intended to be executed (although in Coq, proofs inhabit a lambda calculus and could be “run” in a sense, they usually are just checked). How does BDI relate?

- **Different Goals:** Coq/Lean aim at proving mathematical correctness with absolute guarantees, and extraction is used to get executable code (which then is separate and must be trusted after extraction). BDI's goal is to integrate proving and executing. In BDI, a proof of a function's correctness could *be part of the function's implementation*. This means that BDI is not as rigorous as Coq's logic (unless we embed that logic in BDI) – BDI by itself is not a theorem prover, it's a substrate that can carry proofs from theorem provers.
- **Unified Artifact:** With Coq, you have one artifact that is the proof (in Gallina language) and another that is the program (extracted OCaml or C, for instance). With BDI, it's one artifact: the BDI graph

contains both program and proof. This is closer to the concept of *proof-carrying code* [1] as repeated, where the host need not run a separate proof checker like Coq's kernel; a simpler checker (or the runtime itself) can verify the proof conditions relevant for execution.

- **Expressiveness:** Coq's proofs can express very high-level logical statements (thanks to higher-order logic). BDI's proof metadata could, in theory, encapsulate anything that Coq can express (maybe as a proof term attached), but the BDI runtime isn't expected to verify an arbitrary Coq proof – it might rely on a certificate like a hash that was checked offline, or limit itself to certain proof domains (like linear integer arithmetic proofs or model-checking results that are easy to verify). So while BDI can carry Coq proofs, the actual *checking* might not be as powerful as Coq's typechecker. In practice, one could imagine a BDI node that says “assume verified by Coq” with a ProofTag linking to a Coq certificate. The trust then shifts to whoever checked that certificate (possibly an offline process).
- **Lean's approach to code+proof:** Lean (and projects like Microsoft's Lean-based verified AI) also explore mixing code and proofs. BDI can be seen as an execution platform for such mix: instead of verifying at the level of Lean's high-level language, one would run the program in BDI and check proof obligations in real time. It's a different cut: Coq/Lean prove once, ahead of time; BDI potentially checks continuously, at runtime (or at load-time).

In sum, BDI isn't competing with Coq/Lean but rather providing a way to *deploy* their results in live systems. It unifies the formal proof world with the systems execution world. It's conceivable that in the future, one writes a program in a language, proves properties in Coq, then compiles the whole thing into BDI – resulting in a single deliverable that a user can run with confidence it has proofs built-in. This is more dynamic than traditional proof certificates which are only consulted by humans or offline verifiers.

4.6.5 BDI vs. Hardware Description Languages (Verilog/VHDL) and HDLs/IRs for Hardware (e.g., RTL, SPIR-V): BDI's graph structure might remind one of digital circuit graphs or hardware netlists, and indeed there are parallels:

- **BDI as a Hardware Netlist:** A digital circuit is a graph of logic gates (nodes) with wires (edges). BDI can describe computations in a similar graph form. One could conceptually use BDI to describe hardware at a higher level (especially since BDI can model concurrency and parallel dataflow well). However, BDI nodes are more coarse-grained than logic gates typically; they might represent entire arithmetic operations or even algorithmic steps, whereas an HDL netlist is down to individual flips-flops and gates.
- **Temporal Behavior:** Hardware descriptions often have inherent parallelism and a notion of time (clock cycles). BDI's default semantics are more similar to a software execution (sequential consistency and explicit control flow). To model hardware precisely, one would need to incorporate timing (perhaps via special nodes representing clock ticks or using the execution properties to mark things as combinatorial vs sequential). It's not the primary target of BDI to replace HDLs, but interestingly, by choosing suitable primitives, one could embed an HDL's semantics in BDI (indeed, Chapter 7 discusses interfacing BDI with silicon).
- **SPIR-V and GPU Shaders:** SPIR-V [4] is a binary IR for GPUs, which is also a sort of graph-based SSA (in that it has an SSA form and is stored in binary). SPIR-V is limited to representing computations for GPU shaders and kernels, with specific rules and capabilities (like certain types, no arbitrary pointer

arithmetic beyond what's allowed). BDI can be seen as a generalized superset: anything you can write in SPIR-V (which is mostly numeric kernel code) you can represent in BDI, but not vice versa (SPIR-V wouldn't know what to do with a proof node or an adaptive self-modifying structure). When it comes to optimizing for hardware like GPUs, BDI's region and hint system plays a role analogous to what SPIR-V's execution model does (workgroups, storage classes in SPIR-V correspond to memory regions and hardware units in BDI).

- **Co-Design:** BDI suggests a future where software and hardware might be co-designed more tightly. By having a substrate that can describe computation at a high level but still execute on hardware efficiently, one could feed BDI graphs to synthesis tools to create specialized hardware. For example, a BDI graph for an encryption algorithm, with all its proofs (like proofs of certain invariants), could potentially be synthesized into a cryptographic hardware module where those invariants translate to hardware assertions or formal verification checks on the design. Traditional HDLs don't carry proofs; hardware verification is a separate process. BDI could unify that – the same graph that serves as the source for hardware implementation can contain the assertions to be verified by an SMT solver or model checker.

Therefore, BDI complements HDLs by being at a higher-level (algorithmic), while still capable of dropping down into bit-level logic if needed (with appropriate nodes). It also shares the declarative, parallel nature of hardware descriptions more than sequential programming languages do.

4.6.6 BDI vs. Machine Learning Graph Compilers (XLA, TIRAMISU, etc.): In recent years, systems like Google's XLA [9] and MIT's Tiramisu [10] have been developed to optimize computational graphs (especially for machine learning workloads). They take as input a computational graph (e.g., a TensorFlow or PyTorch graph for XLA, or a scheduling description for Tiramisu) and perform transformations for performance (like fusion, loop tiling, etc.). How does BDI relate?

- **Internal Graph vs. Universal Graph:** XLA operates on a graph of ops (like matrix multiply, convolution, etc.) specifically geared towards ML. Tiramisu represents computations and allows the user to define transformations (schedules) to optimize them, particularly for multi-core and GPU execution. These are domain-specific graph IRs. BDI, on the other hand, is a *universal* graph IR. In fact, one could lower an XLA graph or a Tiramisu representation into BDI (likely losing some domain-specific abstractions but keeping the core computation). Conversely, BDI could serve as the *common language* in which such domain-specific compilers express their result. For example, after Tiramisu computes an optimized schedule for an ML algorithm, it could emit a BDI graph that realizes that schedule with explicit region assignments (like one region per GPU memory tile, etc., captured in RegionIDs) ¹⁶ ¹⁷ .
- **Optimization Capabilities:** Both XLA and Tiramisu focus on performance optimization (fusing operations to reduce memory traffic, reordering loops for locality, vectorization). BDI can incorporate those optimizations as graph transformations (as shown in Figure 4.5 with FMA fusion). The advantage BDI has is that it can leverage semantic information and proofs while optimizing. For instance, XLA has to be careful to maintain numerically equivalent results when fusing operations; BDI might know from a proof that certain transformations won't change the result beyond an acceptable tolerance. Tiramisu requires the user to provide a correct schedule; BDI could potentially explore schedules autonomously and verify if they respect dependencies (since dependencies are explicit, violating them would be obvious or provably incorrect).

- **Runtime Adaptation:** XLA compiles a static graph into optimized code. BDI could allow the graph to adapt at runtime. Consider an ML model that learns a new substructure; BDI can modify the graph accordingly without leaving the execution environment. This is beyond XLA's scope, which is an ahead-of-time compiler. It edges into areas like neural network architectures that evolve (autoML) – BDI would naturally accommodate that because the network is just a BDI graph that can be mutated.
- **Target Hardware:** XLA's output is typically lowered to LLVM IR for CPU or GPU-specific code, or to CUDA for GPUs, etc. Tiramisu likewise generates C++ or low-level code. BDI, in contrast, could target any hardware via its HAL, or even run interpretively. So BDI could actually serve as a back-end for these compilers: rather than generating C++/CUDA, they could generate BDI and let the BDI toolchain handle the final mile to binary code for each platform. The benefit would be that the BDI form can be used for further cross-domain optimizations (maybe combining an ML model graph with a traditional algorithmic piece all in one BDI program, which neither XLA nor Tiramisu alone could do easily).

In conclusion, BDI can be seen as **encompassing** many of these systems' functionalities: like an IR, it can be a compiler target (4.6.1); like a VM, it executes portable code safely (4.6.2, 4.6.3); like a proof framework, it carries verifications (4.6.4); like an HDL, it captures structural computation (4.6.5); and like an ML graph compiler, it enables high-level optimizations (4.6.6). However, by aiming to do all of this in one unified substrate, BDI inevitably introduces complexity – both conceptual and engineering-wise. The hope is that the payoff (a system where **correctness and efficiency** are not at odds but rather achieved together) is worth that complexity.

4.7 Toward a Unified Computing Paradigm

The Binary Decomposition Interface is more than just a new IR or a VM – it represents a shift in how we think about programs and computations. By unifying the layers of abstraction, BDI asks us to imagine a future where *the same representation* underlies design, verification, and execution:

- A scientist formulates a hypothesis as an algorithm, proves its key properties, and compiles it – the result is a BDI graph where every step of the algorithm and every step of the proof are intertwined, ready to be executed on a machine that understands both.
- An operating system is no longer a black box binary with maybe some external proofs of security; instead, it is a living BDI graph that can be inspected for compliance with invariants at any moment, and that can even adapt its strategy (say, scheduling or memory management policy) on the fly, while proving that the adaptation won't violate safety or liveness.
- Hardware and software blur: a critical inner loop might be automatically offloaded to a reconfigurable hardware unit, but not by explicit separate HDL coding – rather, the BDI runtime recognizes a subgraph and maps it to an FPGA at runtime (using partial reconfiguration), all while ensuring the overall system remains coherent and verified.

These scenarios illustrate the philosophical advancement BDI is aiming for. It is **binary-grounded** – all things reduce to bits in the end – but **semantically lifted** – those bits carry meaning accessible at runtime. This could transform how we build intelligent systems: instead of optimizing after the fact or verifying after the fact, the intelligence (learning, reasoning) and assurance (verification) are built in from the start, on the same substrate.

Certainly, realizing this vision poses challenges. The ecosystem needed (from graph construction tools and DSL compilers, to efficient BDIVM implementations, to libraries of proven components) is extensive. Performance concerns will require ingenuity to ensure the richer semantics don't incur overhead when not needed (perhaps via aggressive ahead-of-time compilation that strips out or simplifies checks once proofs are validated). Widespread adoption would require demonstrating that BDI can achieve competitive speeds and that the development productivity is high (the dream being that proving properties as you program in BDI-like frameworks becomes second nature and well-supported by tools, much as testing and debugging are today).

In the following chapters, we delve into specific components that make up the BDI ecosystem. Chapter 5 will explore the design of the BDI Virtual Machine and the execution engine in depth, including how it schedules graphs and checks proofs efficiently. Chapter 6 will discuss the software toolchain – how one writes or generates BDI graphs, including possible programming models. Chapter 7, as previewed, looks at mapping BDI to actual hardware and how hardware might evolve to better support BDI execution. Finally, in Chapter 10, we return to higher-level implications: how BDI and its operating environment (BDIOS) redefine concepts like computational entropy and lay the groundwork for verifiable, adaptive intelligent systems.

Conclusion: The Binary Decomposition Interface stakes out an audacious position: that one representation can simultaneously serve the needs of high-level expressivity, machine-level efficiency, and formal verifiability. It builds on a lineage of ideas ([1] proof-carrying code, [2] typed assembly, [5] safe portable code, [9][10] computation graphs, etc.) and synthesizes them into a single framework. BDI's graph is a place where *logic meets hardware*, where a mathematical proof might sit one edge away from a pointer dereference, all in binary harmony. If successful, BDI could make software fundamentally more transparent and trustworthy – every bit in a running program would, in principle, have a rationale that can be examined or proven. Achieving this unified substrate is a grand challenge, but the path laid out through the rest of this book will, we hope, convince the reader that BDI is not only theoretically intriguing but practically within reach, and potentially transformative for the future of computing and intelligent systems.

References:

- [1] G. C. Necula, "Proof-Carrying Code," *Proc. POPL*, 1997.
- [2] G. Morrisett, D. Walker, K. Crary, and N. Glew, "From System F to Typed Assembly Language," *ACM TOPLAS*, 1999.
- [3] C. Lattner, M. Amini, et al., "MLIR: A Compiler Infrastructure for the End of Moore's Law," arXiv: 2002.11054, 2020.
- [4] Khronos Group, "SPIR-V Specification, Version 1.6," 2023.
- [5] A. Haas *et al.*, "Bringing the Web Up to Speed with WebAssembly," *Proc. PLDI*, 2017.
- [6] X. Leroy, "Formal Verification of a Realistic Compiler," *Commun. ACM*, 52(7):107–115, 2009.
- [7] M. Penninckx, Y. Jugé, F. Piessens, and B. De Sutter, "Proof-Carrying Code for the LLVM Bitcode," *J. Autom. Reason.*, 63(4):929–961, 2019.
- [8] C. Lattner, "Why SSA Isn't Enough," *Chris Lattner's Blog*, 2017.
- [9] Google, "XLA: Optimizing Compiler for Machine Learning," 2022. ([web page](#))
- [10] S. Baghdadi *et al.*, "TIRAMISU: A Polyhedral Compiler for Expressing Fast and Portable Code," *Proc. CC*, 2019.

1 2 8 9 10 11 12 13 14 15 2025-05-05-The Binary Decomposition Interface.pdf

file:///file-U7aYBvUwrBL2chshZnFUzf

3 4 5 6 7 16 17 2025-04-27-Binary Mathematics and Intelligent Systems by Tariq Mohammed
Chapters 1-10.pdf

file:///file-UBwwFk7FpYVMod3u9msGdJ