

Administrivia


- Last project due Friday
- Final Exam
 - Wednesday, December 9th, 3:30-6:30pm
 - Open notes (except textbook), covers all 19 lectures (including topics already on the midterm)
- Final review session Friday (recorded)
 - Bring questions on lecture material
- Extra office hours next week
 - Reload class home page for details

Outline

- 1 Confining code with legacy OSES
- 2 Virtual machines
- 3 Implementing virtual machines
- 4 Binary translation
- 5 Hardware-assisted virtualization
- 6 Memory management optimizations

Confining code with legacy OSes

- Often want to confine code on legacy OSes
- Analogy: Firewalls



- Your machine runs hopelessly insecure software
 - Can't fix it—no source or too complicated
 - *Can* reason about network traffic
- Can we similarly block untrusted code *within* a machine
 - Have OS limit what the code can interact with

Using chroot

- `chroot (char *dir)` “changes root directory”
 - Kernel stores root directory of each process
 - File name “/” now refers to `dir`
 - Accessing “..” in `dir` now returns `dir`
- Need root privileges to call `chroot`
 - But subsequently can drop privileges
- Ideally “Chrooted process” wouldn’t affect parts of the system outside of `dir`
 - Even process still running as root shouldn’t escape `chroot`
- In reality, many ways to cause damage outside `dir`

Escaping chroot

- Re-chroot to a lower directory, then chroot ../.../...
 - Each process has one root directory in process structure
 - Implementation special-cases / (always) & .. in root directory
 - chroot does not always change current directory
 - So chrooting to a lower directory puts you above your new root (Can re-chroot to real system root)
- Create devices that let you access raw disk
- Send signals to or ptrace non-chrooted processes
- Create setuid program for non-chrooted processes to run
- Bind privileged ports, mess with clock, reboot, etc.
- Problem: chroot was not originally intended for security
 - FreeBSD jail, Linux vserver have tried to address problems

System call interposition

- Why not use *ptrace* or other debugging facilities to control untrusted programs?
- Almost any “damage” must result from system call
 - delete files → unlink
 - overwrite files → open/write
 - attack over network → socket/bind/connect/send/recv
 - leak private data → open/read/socket/connect/write ...
- So enforce policy by allowing/disallowing each syscall
 - Theoretically much more fine-grained than chroot
 - Plus don't need to be root to do it
- Q: Why is this not a panacea?


Limitations of syscall interposition

- Hard to know exact implications of a system call
 - Too much context not available outside of kernel (e.g., what does this file descriptor number mean?)
 - Context-dependent (e.g., `/proc/self/cwd`)
- Indirect paths to resources
 - File descriptor passing, core dumps, “unhelpful processes”
- Race conditions
 - Remember difficulty of eliminating TOCCTOU bugs?
 - Now imagine malicious application deliberately doing this
 - Symlinks, directory renames (so “..” changes), ...
- See [\[Garfinkel\]](#) for a more detailed discussion

Outline


- 1 Confining code with legacy OSES
- 2 Virtual machines
- 3 Implementing virtual machines
- 4 Binary translation
- 5 Hardware-assisted virtualization
- 6 Memory management optimizations

Review: What is an OS



- OS is software between applications and hardware/external reality
 - Abstracts hardware to makes applications portable
 - Makes finite resources (memory, # CPU cores) appear much larger
 - Protects processes and users from one another

What if...



- The process abstraction looked just like hardware?

How do process abstraction & HW differ?

Process

Non-privileged registers and instructions

Virtual memory

Errors, signals

File system, directories, files, raw devices

Hardware

All registers and instructions


Both virtual and physical memory, MMU functions, TLB/page tables, etc.

Trap architecture, interrupts

I/O devices accessed using programmed I/O, DMA, interrupts

Virtual Machine Monitor

- Thin layer of software that virtualizes the hardware
 - Exports a virtual machine abstraction that looks like the hardware



Old idea from the 1960s

- See [\[Goldberg\]](#) from 1974
- IBM VM/370 – A VMM for IBM mainframe
 - Multiplex multiple OS environments on expensive hardware
 - Desirable when few machines around
- Interest died out in the 1980s and 1990s
 - Hardware got cheap
 - Compare Windows NT vs. N DOS machines
- Today, VMs are used everywhere
 - Used to solve different problems (software management)
 - But VMM attributes more relevant now than ever

VMM benefits

- Software compatibility
 - VMMs can run pretty much all software
- Can get low overheads/high performance
 - Near “raw” machine performance for many workloads
 - With tricks can have direct execution on CPU/MMU
- Isolation
 - Seemingly total data isolation between virtual machines
 - Leverage hardware memory protection mechanisms
- Encapsulation
 - Virtual machines are not tied to physical machines
 - Checkpoint/migration

OS backwards compatibility

- Backward compatibility is bane of new OSes
 - Huge effort require to innovate but not break
- Security considerations may make it impossible
 - Choice: Close security hole and break apps or be insecure
- Example: Windows XP is end of life
 - Eventually hardware running WinXP will die
 - What to do with legacy WinXP applications?
 - Not all applications will run on later Windows
 - Given the number of WinXP applications, practically any OS change will break something
 - `if (OS == WinXP) ...`
- Solution: Use a VMM to run both WinXP and Win10
 - Obvious for OS migration as well: Windows → Linux

Logical partitioning of servers

- Run multiple servers on same box (e.g., Amazon EC2)
 - Ability to give away less than one machine
Modern CPUs more powerful than most services need
 - 0.10U rack space machine – less power, cooling, space, etc.
 - Server consolidation trend: N machines \rightarrow 1 real machine
- Isolation of environments
 - Printer server doesn't take down Exchange server
 - Compromise of one VM can't get at data of others¹
- Resource management
 - Provide service-level agreements
- Heterogeneous environments
 - Linux, FreeBSD, Windows, etc.

¹ though in practice not so simple because of side-channel attacks [\[Ristenpart\]](#)

Outline

- 1 Confining code with legacy OSes
- 2 Virtual machines
- 3 Implementing virtual machines
- 4 Binary translation
- 5 Hardware-assisted virtualization
- 6 Memory management optimizations

Complete Machine Simulation

- Simplest VMM approach, used by bochs
- Build a simulation of all the hardware
 - CPU – A loop that fetches each instruction, decodes it, simulates its effect on the machine state
 - Memory – Physical memory is just an array, simulate the MMU on all memory accesses
 - I/O – Simulate I/O devices, programmed I/O, DMA, interrupts
- Problem: Too slow!
 - CPU/Memory – 100x CPU/MMU simulation
 - I/O Device – $< 2\times$ slowdown.
 - $100\times$ slowdown makes it not too useful
- Need faster ways of emulating CPU/MMU

Virtualizing the CPU

- Observations: Most instructions are the same regardless of processor privileged level
 - Example: `incl %eax`
- Why not just give instructions to CPU to execute?
 - One issue: Safety – How to get the CPU back? Or stop it from stepping on us? How about `cli/halt`?
 - Solution: Use protection mechanisms already in CPU
- Run virtual machine's OS directly on CPU in unprivileged user mode
 - “Trap and emulate” approach
 - Most instructions just work
 - Privileged instructions trap into monitor and run simulator on instruction
 - Makes some assumptions about architecture

Virtualizing traps

- What happens when an interrupt or trap occurs
 - Like normal kernels: we trap into the monitor
- What if the interrupt or trap should go to guest OS?
 - Example: Page fault, illegal instruction, system call, interrupt
 - Re-start the guest OS simulating the trap
- x86 example:
 - Give CPU an IDT that vectors back to VMM
 - Look up trap vector in VM's "virtual" IDT
 - Push virtualized %cs, %eip, %eflags, on stack
 - Switch to virtualized privileged mode

Virtualizing memory

- Basic MMU functionality:
 - OS manages physical memory (0...MAX_MEM)
 - OS sets up page tables mapping VA \rightarrow PA
 - CPU accesses to VA should go to PA (if paging off, PA = VA)
 - Used for every instruction fetch, load, or store
- Need to implement a virtual “physical memory”
 - Logically need additional level of indirection
 - VM's *Guest VA* \rightarrow VM's *Guest PA* \rightarrow *Host PA*
 - Note “Guest physical” memory no longer means hardware bits
 - Hardware is host physical memory (a.k.a. machine memory)
- Trick: Use hardware MMU to simulate virtual MMU
 - Point hardware at *shadow page table*
 - Directly maps Guest VA \rightarrow Host PA

Memory mapping summary



physical machine

virtual machine




Shadow page tables

- VMM responsible for maintaining *shadow* PT
 - And for maintaining its consistency (including TLB flushes)
- Shadow page tables are a cache
 - Have *true page faults* when page not in VM's guest page table
 - Have *hidden page faults* when just misses in shadow page table
- On a page fault, VMM must:
 - Lookup guest VPN → guest PPN in guest's page table
 - Determine where guest PPN is in host physical memory
 - Insert guest VPN → host PPN mapping in shadow page table
 - Note: Monitor can demand-page the virtual machine
- Uses hardware protection

Shadow PT issues

- Hardware only ever sees shadow page table
 - Guest OS only sees it's own VM page table, never shadow PT
- Consider the following
 - Guest OS has a page table T mapping $V_U \rightarrow P_U$
 - T itself resides at guest physical address P_T
 - Another guest page table entry maps $V_T \rightarrow P_T$
 - VMM stores P_U in host physical address M_U and P_T in M_T
- What can VMM put in shadow page table?
 - Safe to map $V_T \rightarrow M_T$ *or* $V_U \rightarrow M_U$
- **Not safe to map both simultaneously!**
 - If OS writes to P_T , may make $V_U \rightarrow M_U$ in shadow PT incorrect
 - If OS reads/writes V_U , may require accessed/dirty bits to be changed in P_T (hardware can only change shadow PT)

Illustration



- Option 1: Page table accessible at V_T , but changes won't be reflected in shadow PT or TLB; access to V_U dangerous
- Option 2: V_U accessible, but hardware sets accessed/dirty bits only in shadow PT, not in guest PT at P_T/M_T

Tracing

- VMM needs to get control on some memory accesses
- Guest OS changes previously used mapping in VM PT
 - Must intercept to invalidate stale mappings in shadow PT, TLB
 - Note: OS *should* use `invlpg` instruction, which would trap to VMM – but in practice many/most OSes are sloppy about this
- Guest OS accesses page when its VM PT is accessible
 - Accessed/dirty bits in VM PT may no longer be correct
 - Must intercept to fix up VM PT (or make VM PT inaccessible)
- Solution: *Tracing*
 - To track page access, make VPN(s) invalid in shadow PT
 - If guest OS accesses page, will trap to VMM w. page fault
 - VMM can emulate the result of memory access & restart guest OS, just as an OS restarts a process after a page fault

Tracing vs. hidden faults

- Suppose VMM never allowed access to VM PTs?
 - Every PTE access would incur the cost of a tracing fault
 - Very expensive when OS changes lots of PTEs
- Suppose OS allowed access to *most* page tables (except very recently accessed regions)
 - Now lots of hidden faults when accessing new region
 - Plus overhead to pre-compute accessed/dirty bits from shadow PT as page tables preemptively made valid in shadow PT
- Makes for complex trade-offs
 - But adaptive binary translation (later) can make this better

I/O device virtualization

- Types of communication
 - Special instruction – `in/out`
 - Memory-mapped I/O (PIO)
 - Interrupts
 - DMA
- Make `in/out` and PIO trap into monitor
- Use tracing for memory-mapped I/O
- Run simulation of I/O device
 - Interrupt – Tell CPU simulator to generate interrupt
 - DMA – Copy data to/from physical memory of virtual machine

CPU virtualization requirements

- Need protection levels to run VMs and monitors
- All unsafe/privileged operations should trap
 - Example: disable interrupt, access I/O dev, ...
 - x86 problem: `popfl` (different semantics in different rings)
- Privilege level should not be visible to software
 - Software shouldn't be able to query and find out it's in a VM
 - x86 problem: `movw %cs, %ax`
- Trap should be transparent to software in VM
 - Software in VM shouldn't be able to tell if instruction trapped
 - x86 problem: traps can destroy machine state (E.g., if internal segment register was out of sync with GDT)
- See [\[Goldberg\]](#) for a discussion

Outline

- 1 Confining code with legacy OSes
- 2 Virtual machines
- 3 Implementing virtual machines
- 4 Binary translation
- 5 Hardware-assisted virtualization
- 6 Memory management optimizations

Binary translation

- Cannot directly execute guest OS kernel code on x86
 - Can maybe execute most user code directly
 - But how to get good performance on kernel code?
- Original VMware solution: binary translation
 - Don't run slow instruction-by-instruction emulator
 - Instead, translate guest kernel code into code that runs in fully-privileged kernel mode, but acts safely²
- Challenges:
 - Don't know the difference between code and data (guest OS might include self-modifying code)
 - Translated code may not be the same size as original
 - Prevent translated code from messing with VMM memory
 - Performance, performance, performance, . . .

² actually CPL 1, so that the VMM has its own exception stack

VMware binary translator

- VMware translates kernel dynamically (like a JIT)
 - Start at guest `eip`
 - Accumulate up to 12 instructions until next control transfer
 - Translate into binary code that can run in VMM context
- Most instructions translated identically
 - E.g., regular `movl` instructions
- Use segmentation to protect VMM memory
 - VMM located in high virtual addresses
 - Segment registers “truncated” to block access to high VAs
 - `gs` segment not truncated; use it to access VMM data
 - Any guest use of `gs` (rare) can’t be identically translated

Details/examples from [\[Adams & Agesen\]](#)

Control transfer

- All branches/jumps require indirection

- Original:

```
isPrime: mov %edi, %ecx # %ecx = %edi (a)
          mov $2, %esi  # i = 2
          cmp %ecx, %esi # is i >= a?
          jge prime     # jump if yes
          ...
```

- C source:

```
int
isPrime (int a)
{
    for (int i = 2; i < a; i++) {
        if (a % i == 0)
            return 0;
    }
    return 1;
}
```

Control transfer

- All branches/jumps require indirection
- Original:

```
isPrime: mov %edi, %ecx # %ecx = %edi (a)
        mov $2, %esi   # i = 2
        cmp %ecx, %esi  # is i >= a?
        jge prime      # jump if yes
        ...
```
- Translated:

```
isPrime': mov %edi, %ecx # IDENT
          mov $2, %esi
          cmp %ecx, %esi
          jge [takenAddr] # JCC
          jmp [fallthrAddr]
```
- Brackets ([...]) indicate *continuations*
 - First time jumped to, target untranslated; translate on demand
 - Then fix up continuation to branch to translated code
 - Can elide [fallthrAddr] if fallthrough next translated

Non-identically translated code

- PC-relative branches & Direct control flow
 - Just compensate for output address of translator on target
 - Insignificant overhead
- Indirect control flow
 - E.g., jump though register (function pointer) or `ret`
 - Can't assume code is "normal" (e.g., must faithfully `ret` even if stack doesn't have return address)
 - Look up target address in hash table to see if already translated
 - "Single-digit percentage" overhead
- Privileged instructions
 - Appropriately modify VMM state
 - E.g., `cli` \implies `vcpu.flags.IF = 0`
 - Can be faster than original!

Adaptive binary translation

- One remaining source of overhead is tracing faults
 - E.g., when modifying page table or descriptor table
- Idea: Use binary translation to speed up
 - E.g., translate write of PTE into write of guest & shadow PTE
 - Translate read of PTE to get accessed & dirty bits from shadow
- Problem: Which instructions to translate?
- Solution: “innocent until proven guilty” model
 - Initially always translate as much code identically as possible
 - Track number of tracing faults caused by an instruction
 - If high number, re-translate to non-identical code
 - May call out to interpreter, or just jump to new code

Outline

- 1 Confining code with legacy OSes
- 2 Virtual machines
- 3 Implementing virtual machines
- 4 Binary translation
- 5 Hardware-assisted virtualization
- 6 Memory management optimizations

Hardware-assisted virtualization

- Both Intel and AMD now have hardware support
 - Different mechanisms, similar concepts
 - This lecture covers AMD (see [\[AMD Vol 2\]](#), Ch. 15)
 - For Intel details, see [\[Intel Vol 3c\]](#)
- VM-enabled CPUs support new *guest* mode
 - This is separate from kernel/user modes in bits 0–1 of %cs
 - Less privileged than *host* mode (where VMM runs)
 - Some sensitive instructions trap in guest mode (e.g., load %cr3)
 - Hardware keeps shadow state for many things (e.g., %eflags)
- Enter guest mode with `vmrun` instruction
 - Loads state from hardware-defined 1-KiB VMCB data structure
- Various events cause EXIT back to host mode
 - On EXIT, hardware saves state back to VMCB

VMCB control bits

- *Intercept vector* specifies what ops should cause EXIT
 - One bit for each of %cr0–%cr15 to say trap on read
 - One bit for each of %cr0–%cr15 to say trap on write
 - 32 analogous bits for the debug registers (%dr0–%dr15)
 - 32 bits for whether to intercept exception vectors 0–31
 - Bits for various other events (e.g., NMI, SMI, ...)
 - Bit to intercept writes to sensitive bits of %cr0
 - 8 bits to intercept reads and writes of IDTR, GDTR, LDTR, TR
 - Bits to intercept rdtsc, rdpmc, pushf, popf, vmrun, hlt, invlpg, int, iret, in/out (to selected ports), ...
- EXIT code and reason (e.g., which inst. caused EXIT)
- Other control values
 - Pending virtual interrupt, event/exception injection

Guest state saved in VMCB

- Saved guest state
 - Full segment registers (i.e., base, lim, attr, not just selectors)
 - Full GDTR, LDTR, IDTR, TR
 - Guest %cr3, %cr2, and other cr/dr registers
 - Guest %eip and %eflags (%rip & %rflags for 64-bit processors)
 - Guest %rax register
- Entering/exiting VMM more expensive than syscall
 - Have to save and restore large VM-state structure

Hardware vs. Software virtualization

- HW VM makes implementing VMM much easier
 - Avoids implementing binary translation (BT)
- Hardware VM is better at entering/exiting kernel
 - E.g., Apache on Windows benchmark: one address space, lots of syscalls, hardware VM does better [\[Adams\]](#)
 - Apache on Linux w. many address spaces: lots of context switches, tracing faults, etc., Software faster [\[Adams\]](#)
- Fork with copy-on-write bad for both HW & BT
 - [\[Adams\]](#) reports fork benchmark where BT-based virtualization $37\times$ and HW-based $106\times$ slower than native!
- Today, CPUs support *nested paging*
 - Eliminates shadow PT & tracing faults, simplifies VMM
 - Guests can now manipulate %cr3 w/o VM EXIT
 - But dramatically increases cost of TLB misses

Outline

- 1 Confining code with legacy OSes
- 2 Virtual machines
- 3 Implementing virtual machines
- 4 Binary translation
- 5 Hardware-assisted virtualization
- 6 Memory management optimizations

ESX memory management [Waldspurger]

- Virtual machines see virtualized physical memory
 - Can let VMs use more “physical” memory than in machine
- How to apportion memory between machines?
- VMware ESX has three parameters per VM:
 - **min** – Don't bother running w/o this much machine memory
 - **max** – Amount of guest physical memory VM OS thinks exists
 - **share** – How much memory to give VM relative to other VMs
- Straw man: Allocate based on share, use LRU paging
 - OS already uses LRU \Rightarrow double paging
 - OS will re-cycle whatever “physical” page VMM just paged out
 - So better to do random eviction
- Next: 3 cool memory management tricks

Reclaiming pages

- Normally OS just uses all available memory
 - But some memory much more important than other memory
 - E.g., buffer cache may contain old, clean buffers; OS won't discard if doesn't need memory... but VMM may need memory
- Idea: Have guest OS return memory to VMM
 - Then VMM doesn't have to page memory to disk
- ESX trick: Balloon driver
 - Special pseudo-device driver in supported guest OS kernels
 - Communicates with VMM through special interface
 - When VMM needs memory, allocates many pages in guest OS
 - Balloon driver tells VMM to re-cycle its private pages

Sharing pages across VMs

- Often run many VMs with same OS, programs
 - Will result in many host physical pages containing same data
- Idea: Use 1 host physical page for all copies of guest physical page (in any virtual machine)
- Keep big hash table mapping: Hash(contents)→info
 - If host physical page mapped once, info is VM/PPN where mapped. In that case, Hash is only a hint, as page may have changed
 - If machine page mapped copy-on-write as multiple physical pages, info is just reference count
- Scan OS pages randomly to populate hash table
- Always try sharing a page before paging it out

Idle memory tax

- Need machine page? What VM to take it from?
- Normal proportional share scheme
 - Reclaim from VM with lowest “shares-to-pages” (S/P) ratio
 - If A & B both have $S = 1$, reclaim from larger VM
 - If A has twice B ’s share, can use twice the machine memory
- High-priority VMs might get more memory than needed
- Solution: Idle-memory tax
 - Use statistical sampling to determine a VM’s % idle memory (randomly invalidate pages & count the number faulted back)
 - Instead of S/P , reclaim from VM with lowest $S / (P(f + k(1 - f)))$.
 f = fraction of non-idle pages; k = “idle page cost” parameter.
 - Be conservative & overestimate f to respect priorities
(f is max of slow, fast, and recent memory usage samples)